

Coloured Petri Nets

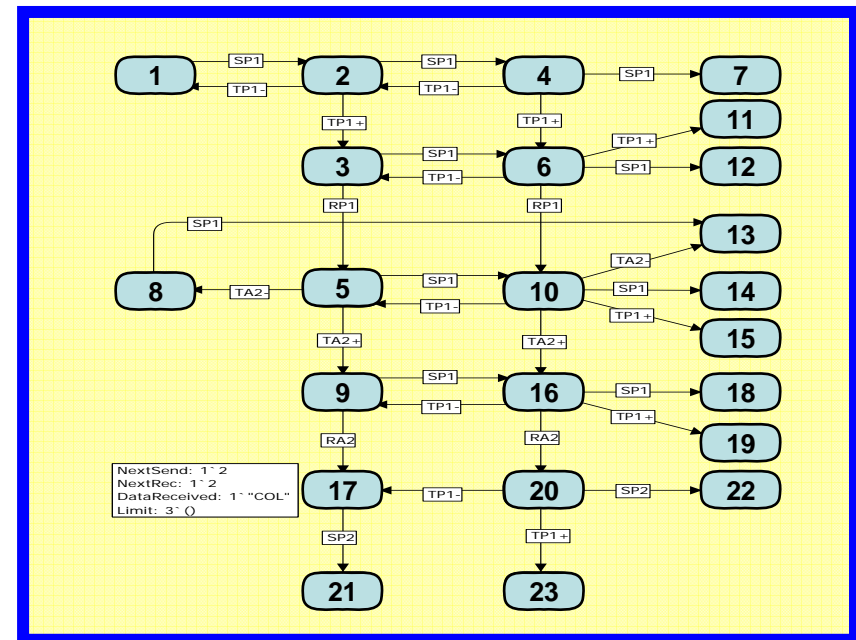
Modelling and Validation of Concurrent Systems

Chapter 7: State Spaces and Behavioural Properties

Kurt Jensen &
Lars Michael Kristensen

{kjensen,lmkristensen}
@daimi.au.dk

© February 2008



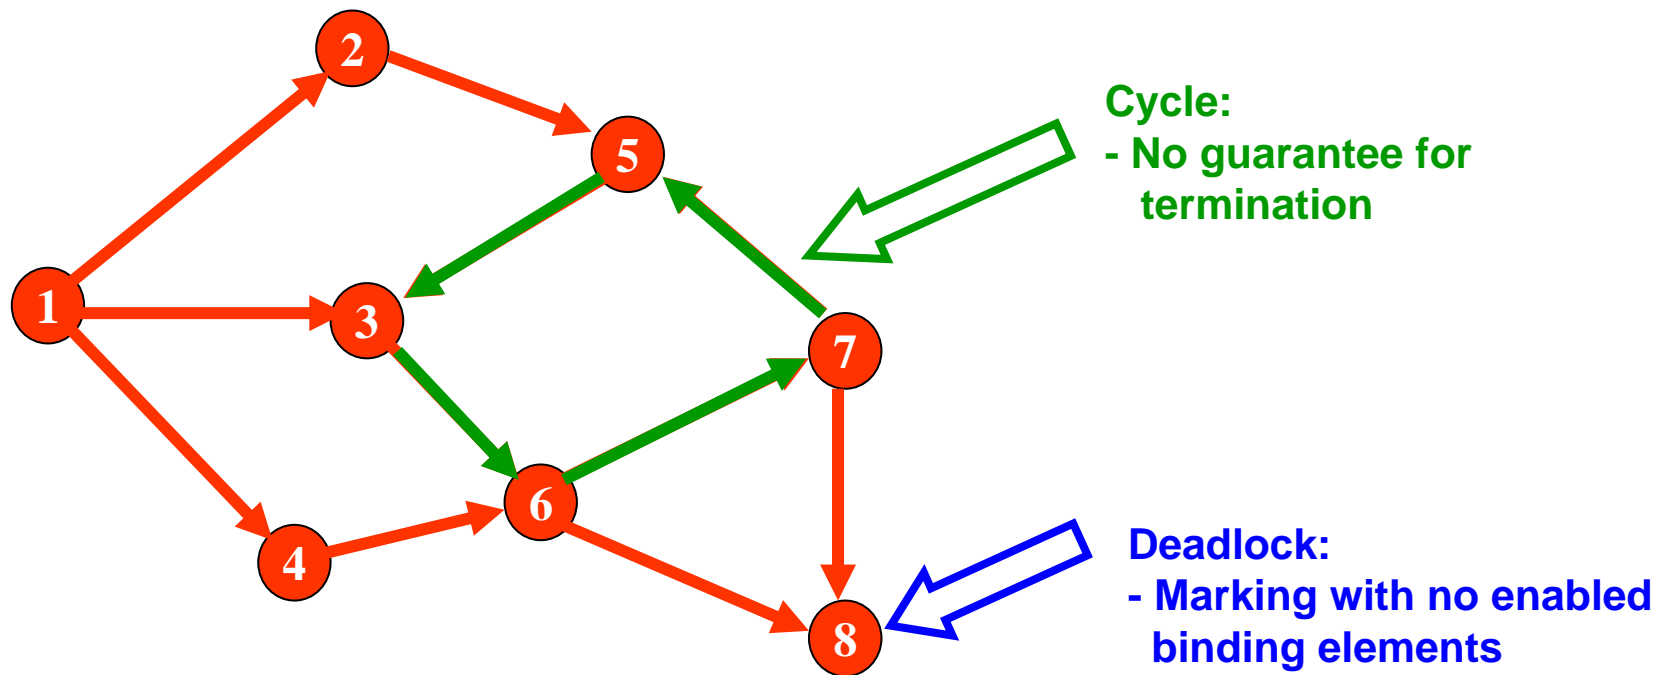
UNIVERSITY OF AARHUS

Modelling and Validation of Distributed Systems Group
Department of Computer Science

Kurt Jensen and Lars M. Kristensen
Coloured Petri Nets

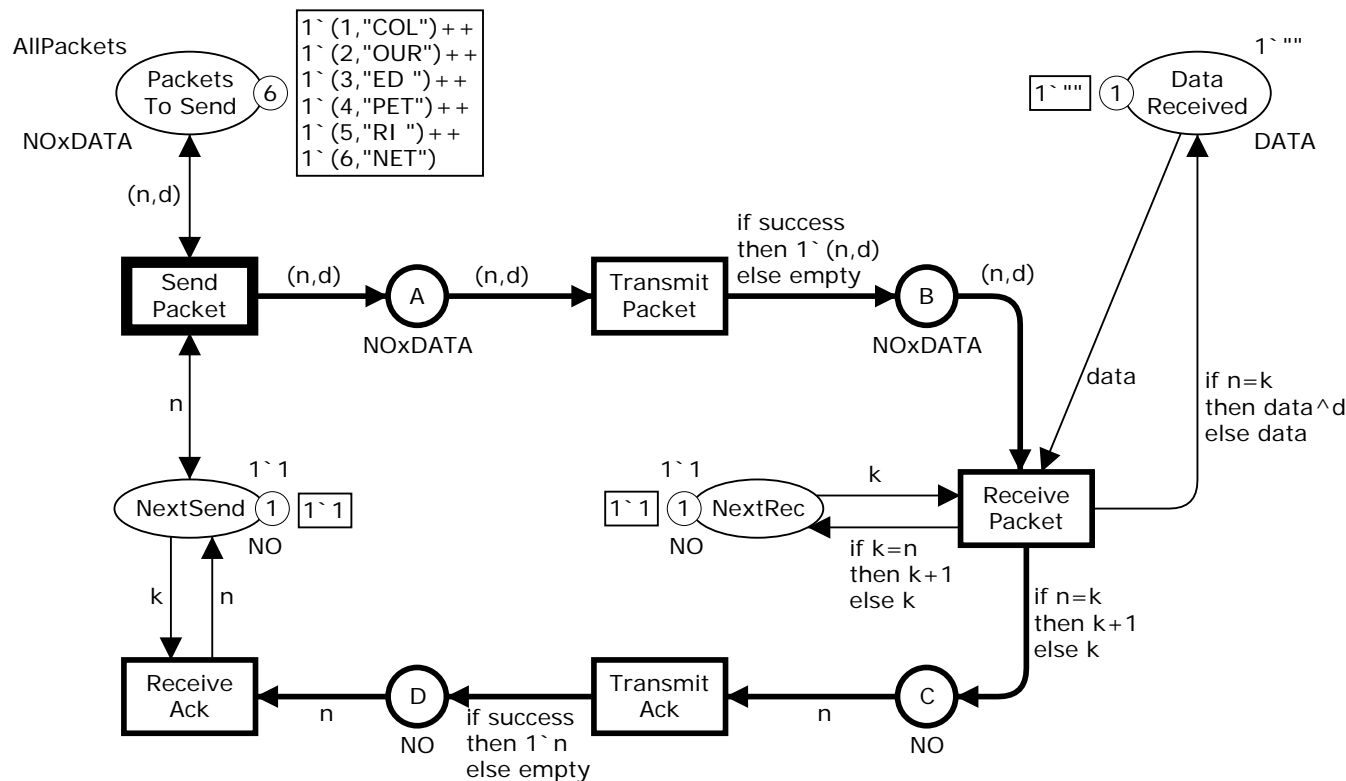
State spaces

- A state space is a **directed graph** with:
 - A **node** for each **reachable marking** (state).
 - An **arc** for each **occurring binding element**.
- **State spaces** can be used to investigate the **behavioural properties** of the CPN model.



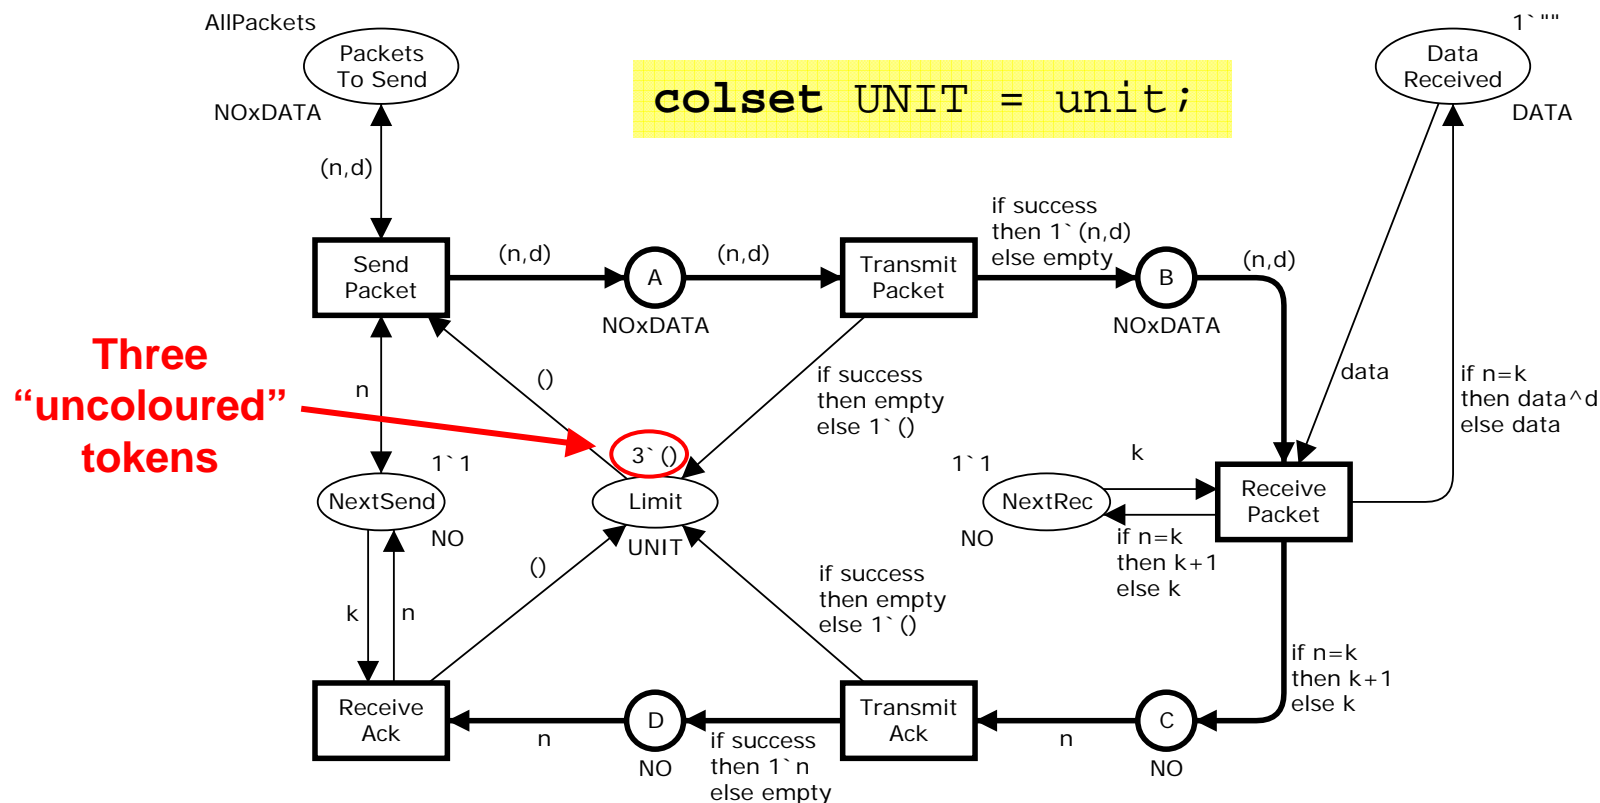
Simple protocol

- **SendPacket** can occur an **unlimited** number of times producing an **unlimited** number of **tokens** on place A.
- This means that the **state space** becomes **infinite**.



Simple protocol for state space analysis

- We add a **new place Limit**, which limits the **total number of tokens** on the **buffer places A, B, C, and D**.
- This makes the **state space finite**.




State space

- Construction of the state space starts with the processing of node 1 which represents the initial marking.
- Node 1 has one enabled binding element:

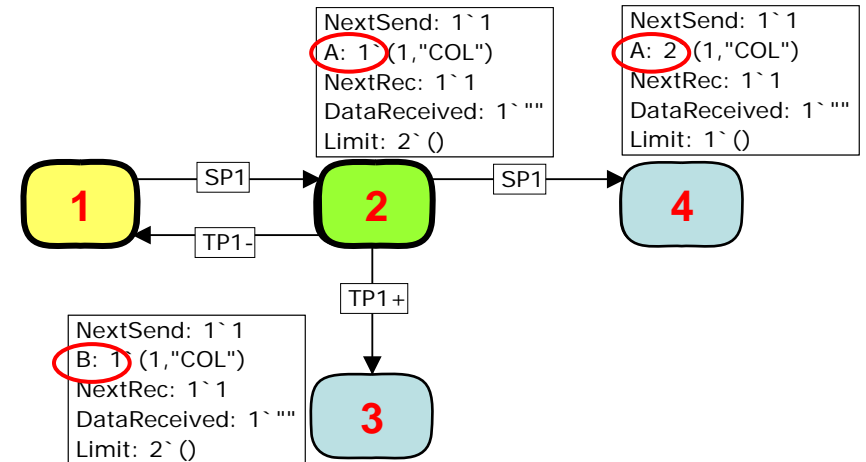
$SP1 = (\text{SendPacket}, \langle n=1, d=\text{"COL"} \rangle)$

- This gives us one new arc and one new node 2.
- Node 2 has one copy of data packet 1 on place A.
- Node 1 is now marked as processed (thick border line).



State space

- Next we **process** node 2.
- It has **three** enabled binding elements:



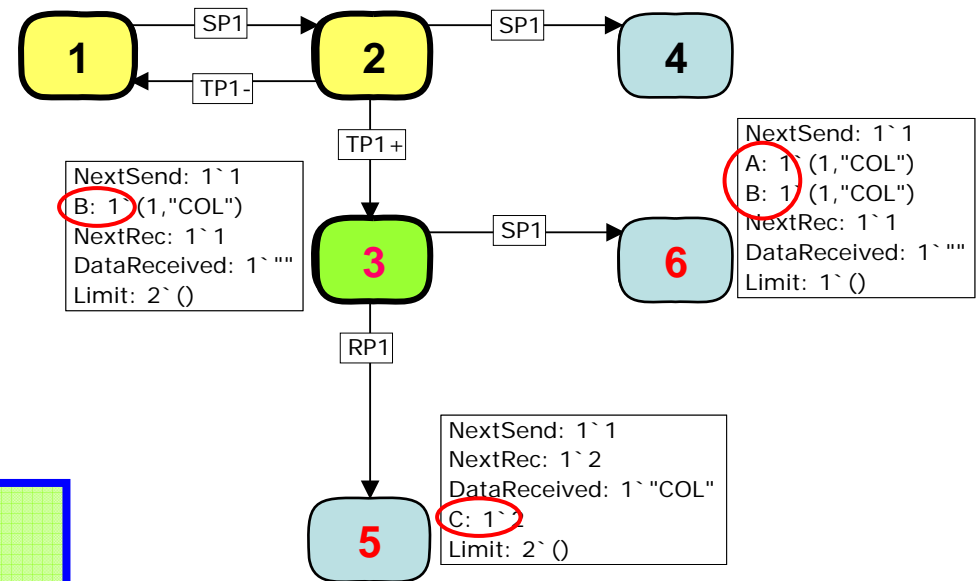
SP1 = (SendPacket, <n=1, d="COL">)
 TP1+ = (TransmitPacket, <n=1, d="COL", success=true>)
 TP1- = (TransmitPacket, <n=1, d="COL", success=false>)

- This gives us **three new arcs** and **two new nodes 3 and 4**.
- Node 3** has one copy of data packet 1 on place B.
- Node 4** has two copies of data packet 1 on place A.
- Node 2** is now marked as **processed** (thick border line).

State space

- Next we **choose** one of the **unprocessed** nodes: 3.
- It has **two** enabled binding elements:

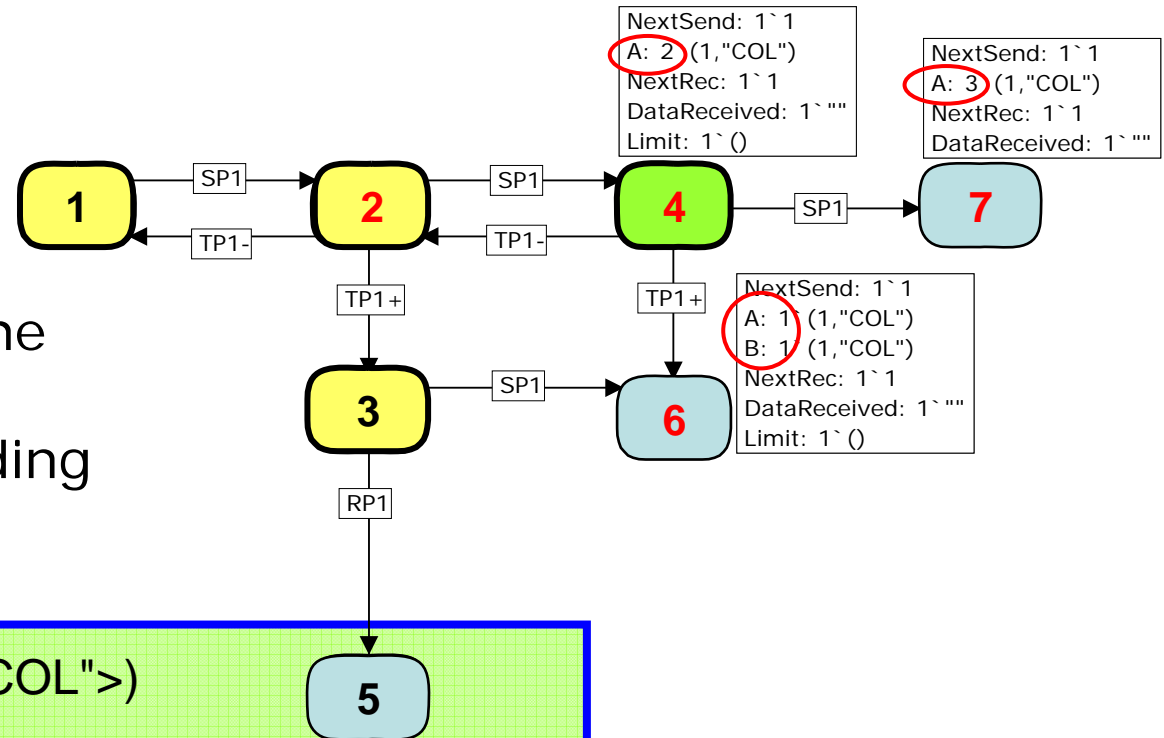
$SP1 = (\text{SendPacket}, \langle n=1, d=\text{"COL"} \rangle)$
 $RP1 = (\text{ReceivePacket}, \langle n=1, d=\text{"COL"}, k=1, data = \text{""} \rangle)$



- This gives us **two new arcs** and **two new nodes 5 and 6**.
- Node 5** has one copy of acknowledgement 2 on place C.
- Node 6** has one copy of packet 1 on place A and another on place B.
- Node 3** is now marked as **processed** (thick border line).

State space

- Next we **choose** one of the **unprocessed** nodes: 4.
- It has **three** enabled binding elements:



SP1 = (SendPacket, <n=1, d="COL">)
 TP1+ = (TransmitPacket, <n=1, d="COL", success=true>)
 TP1- = (TransmitPacket, <n=1, d="COL", success=false>)

- This gives us a **three new arcs** and **one new node** 7.
- Node 7** has three copies of data packet 1 on place A.
- Node 4** is now marked as **processed** (thick border line).

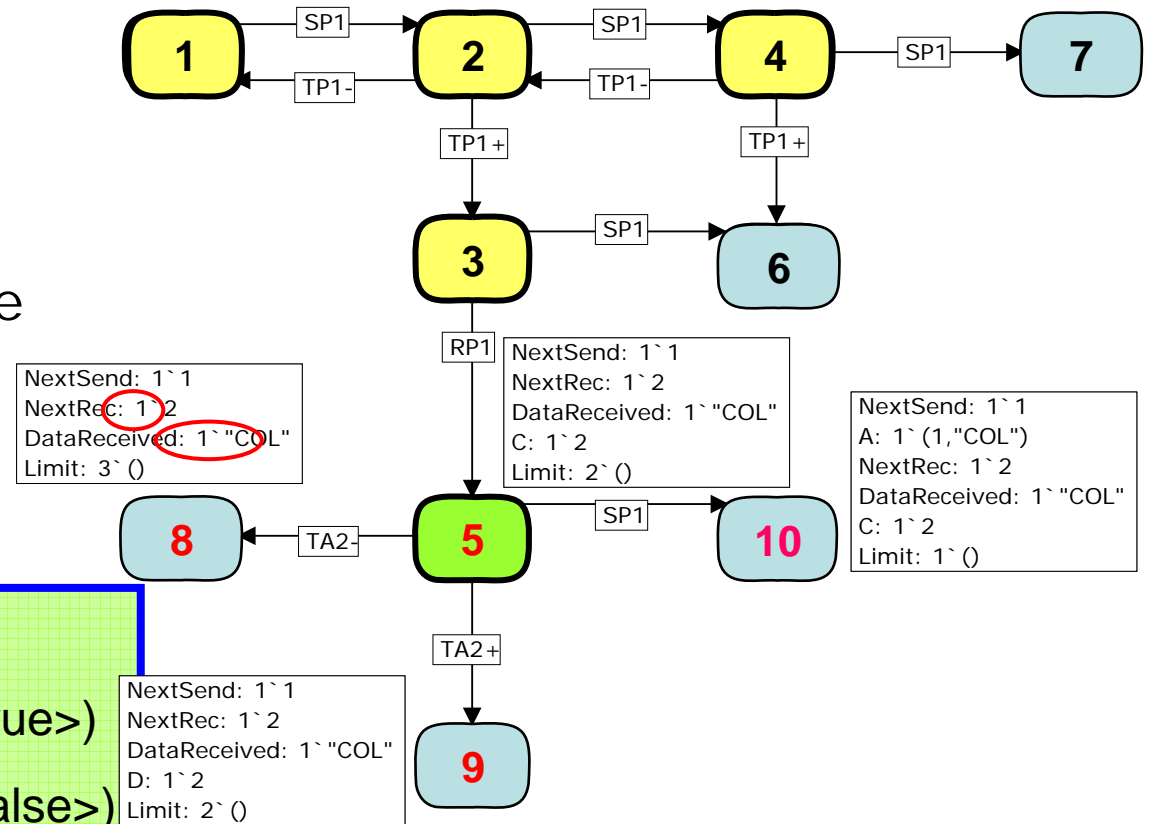
State space

- Next we **choose** one of the **unprocessed** nodes: 5.
- It has **three** enabled binding elements:

SP1 = (SendPacket, <n=1, d="COL">)

TA2+ = (TransmitAck, <n=2, success=true>)

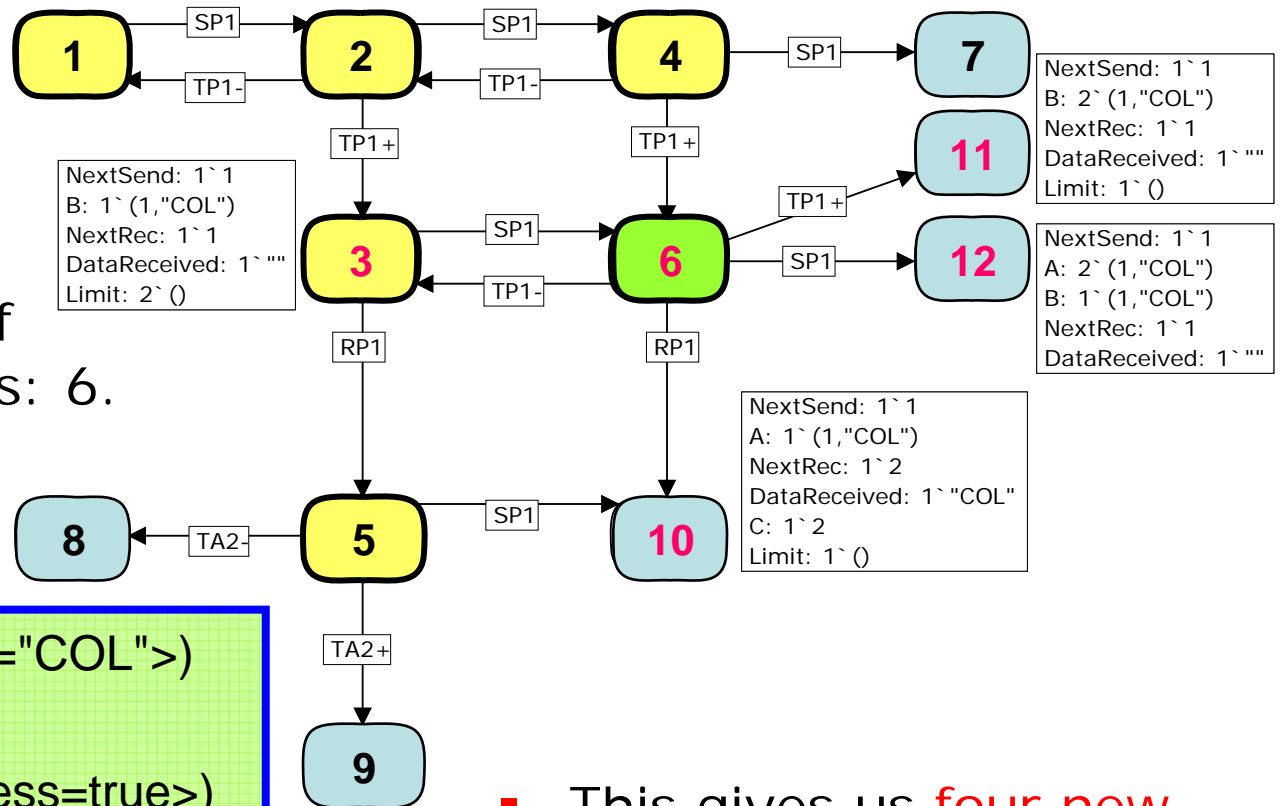
TA2- = (TransmitAck, <n=2, success=false>)



- This gives us **three new arcs** and **three new nodes** 8, 9 and 10.
- Node 8** is identical to the initial marking except that NextRec and Data Received have been changed.
- Node 5** is now marked as **processed** (thick border line).

State space

- Next we **choose** one of the **unprocessed** nodes: 6.
- It has **four** enabled binding elements:



SP1 = (SendPacket, <n=1, d="COL">)

TP1+ = (TransmitPacket, <n=1, d="COL", success=true>)

TP1- = (TransmitPacket, <n=1, d="COL", success=false>)

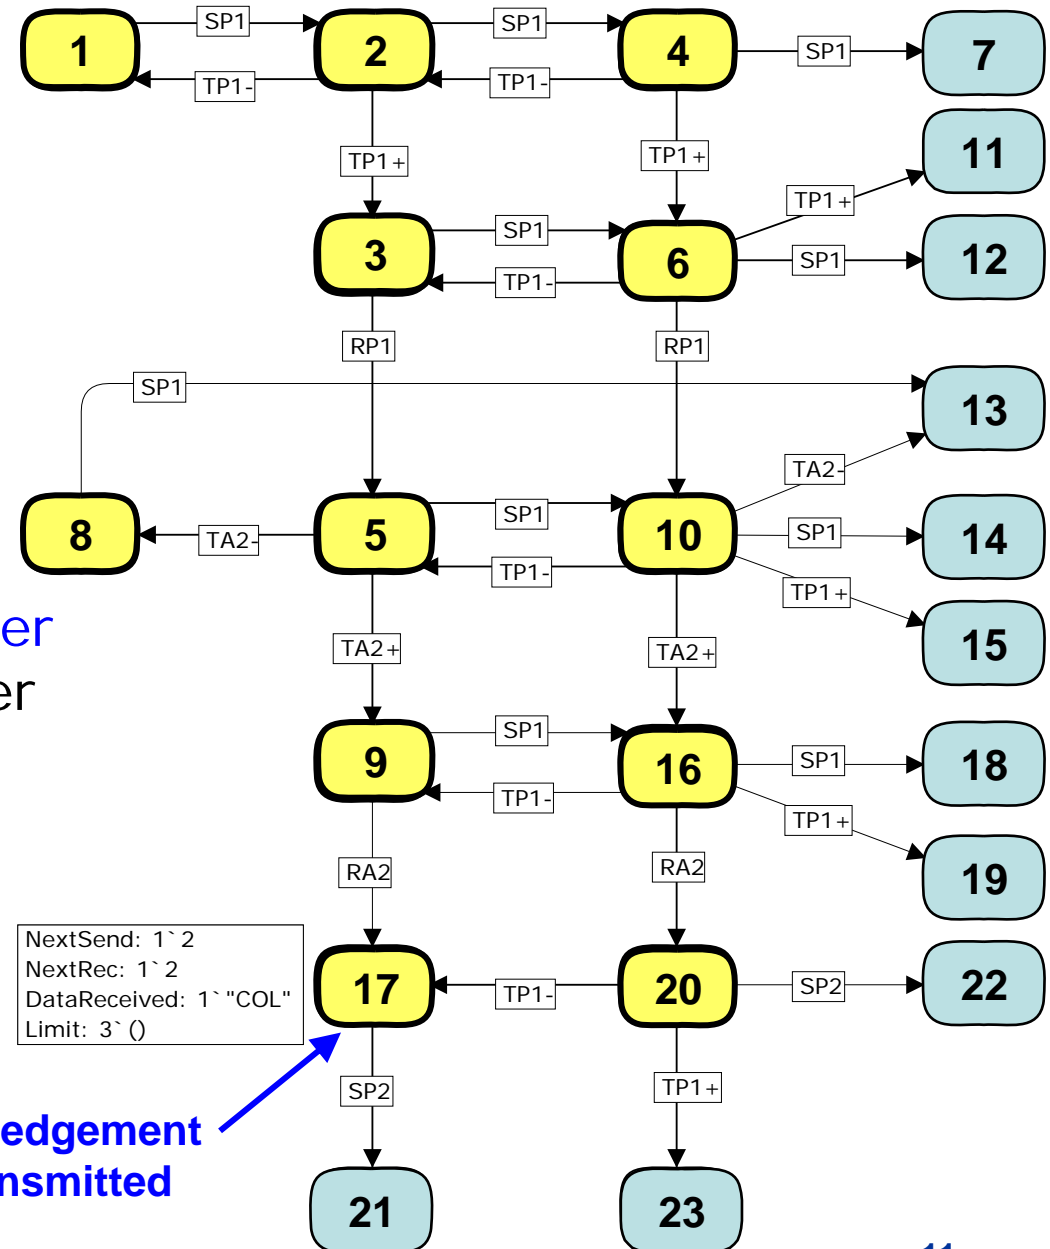
RP1 = (ReceivePacket, <n=1, d="COL", k=1, data = "">)

- This gives us **four new arcs** and **two new nodes** 11 and 12.
- Node 6** is now marked as **processed**.



State space

- We **continue** to **process** the nodes **one by one**.
- If the state space is **finite** construction **terminates** when all **reachable markings** have been **processed**.
- Otherwise, we **continue forever** – obtaining a larger and larger part of the state space.
- This **partial state space** is **visualised** using the drawing facilities of the **CPN state space tool**.




**Packet no. 1 and its acknowledgement
have been successfully transmitted**

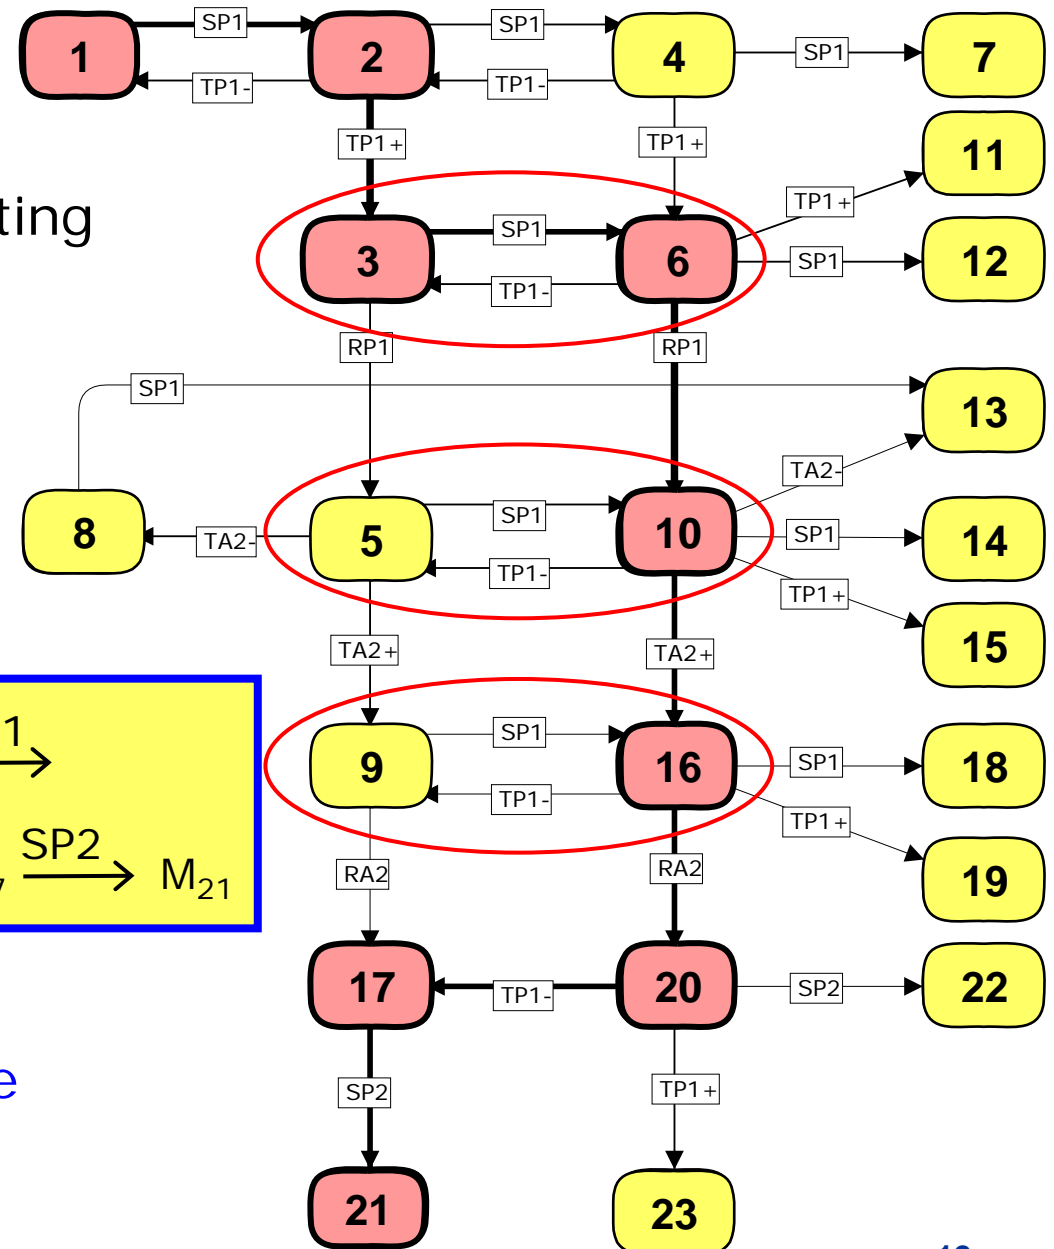


Directed path

- A **directed path** is an alternating sequence of nodes and arcs.
- Each **directed path** in the state space corresponds to an **occurrence sequence** where all steps contain a **single binding element**.




- **Loops** can be **repeated**.
- **Infinite number** of **occurrence sequences**.




Strongly connected components

- A **strongly connected component (SCC)** is a **maximal subgraph** in which **all nodes** are reachable **from each other**.
- The SCCs are **mutually disjoint**.
- Each node is in **exactly one** SCC.
- **SCC graph** contains:
 - A **node** for **each** SCC.
 - An **arc** from S_i to S_j for each state space arc from a node $n_i \in S_i$ to a node $n_j \in S_j$ ($i \neq j$).
- The SCC graph is **acyclic**.



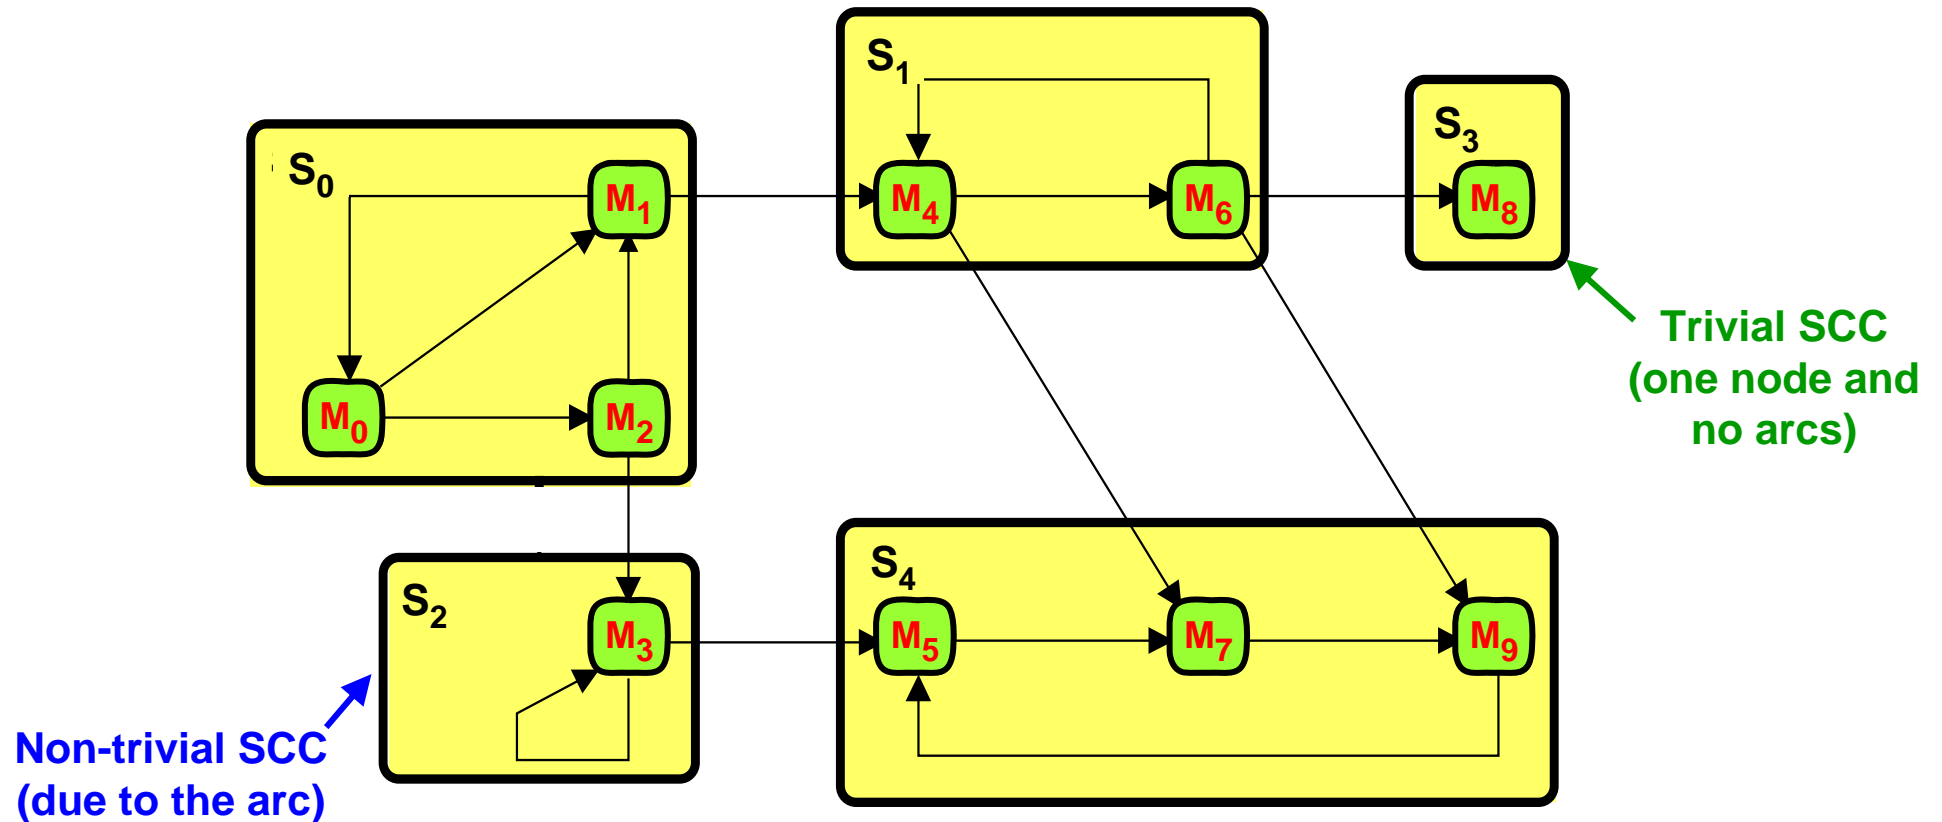
State space (example)

- 10 nodes and 16 arcs.



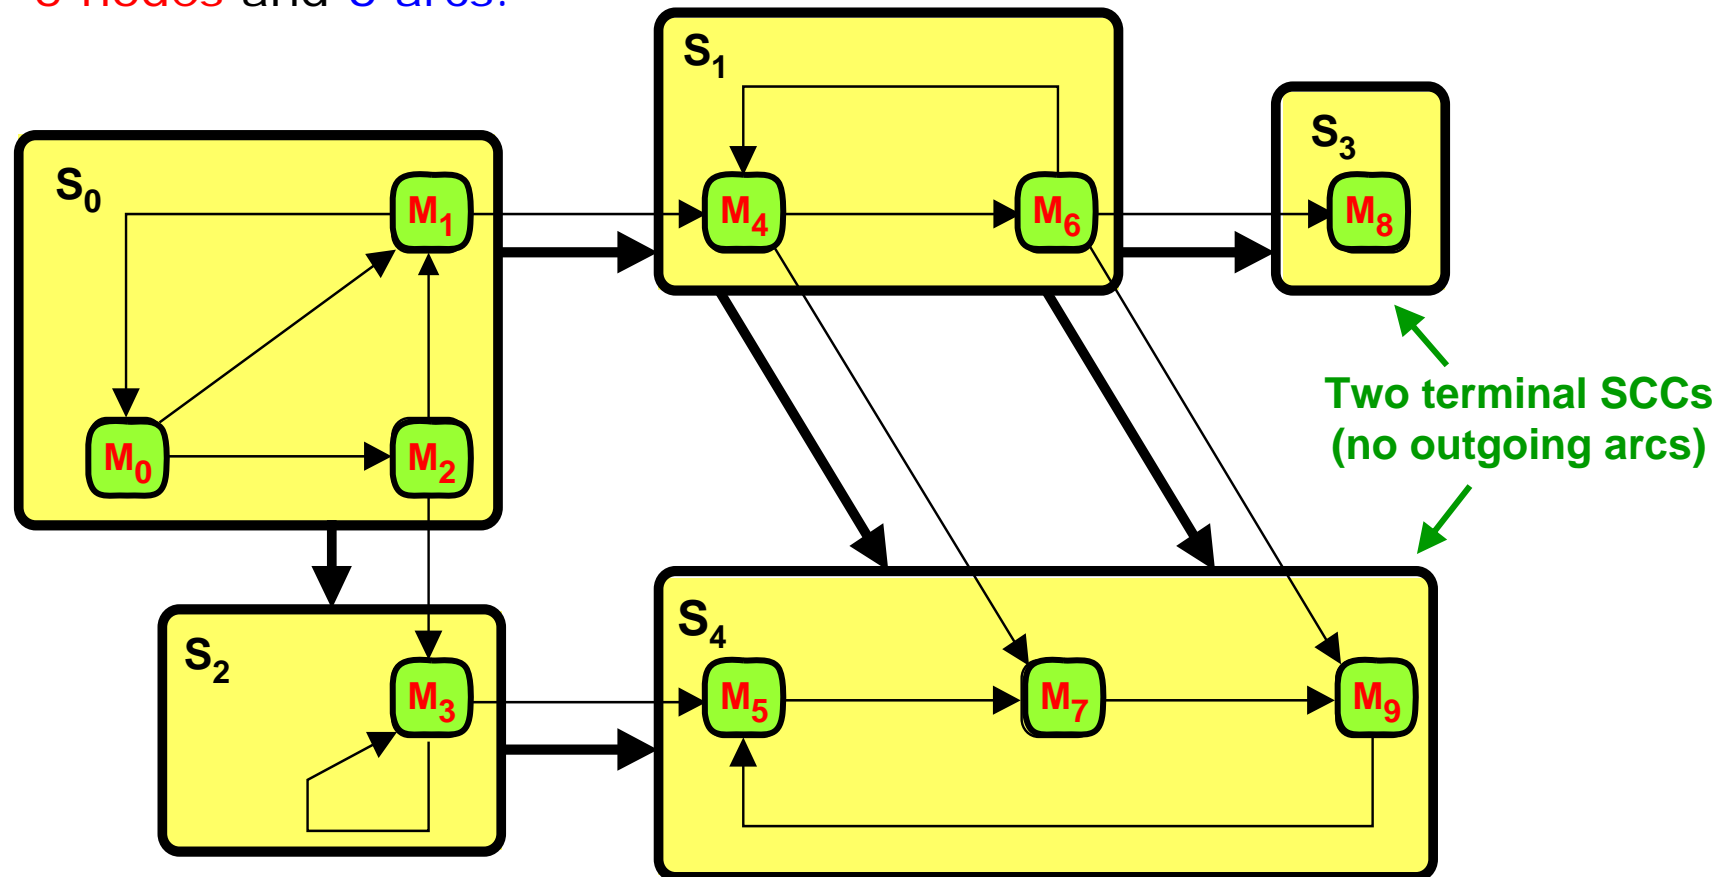
Strongly connected components

- 5 different SCCs.



SCC graph

- 5 nodes and 6 arcs.



State space construction and analysis

- State spaces may be very large and hence we need computer tools to construct and analyse them.
- Analysis of the state space starts with the generation of the state space report.
- This is done totally automatic.
- The report contains a lot of useful information about the behavioural properties of the CPN model.
- The report is excellent for locating errors or increase our confidence in the correctness of the system.



State space report

- The **state space report** contains information about **standard behavioural properties** which make sense for **all CPN models**:
 - **Size** of the state space and the **time** used to generate it.
 - **Bounds** for the **number of tokens** on each place and information about the **possible token colours**.
 - **Home markings**.
 - **Dead markings**.
 - **Dead** and **live** transitions.
 - **Fairness** properties for **transitions**.



State space report: size and time

State Space Statistics

State Space

Nodes: 13.215
Arcs: 52.784
Secs: 53
Status: Full

Scs Graph

Nodes: 5.013
Arcs: 37.312
Secs: 2

- State space contains more than 13.000 nodes and more than 52.000 arcs.
- The state space was constructed in less than one minute and it is full – i.e. contains all reachable markings.
- The SCC graph is smaller. Hence we have cycles.
- The SCC graph was constructed in 2 seconds.



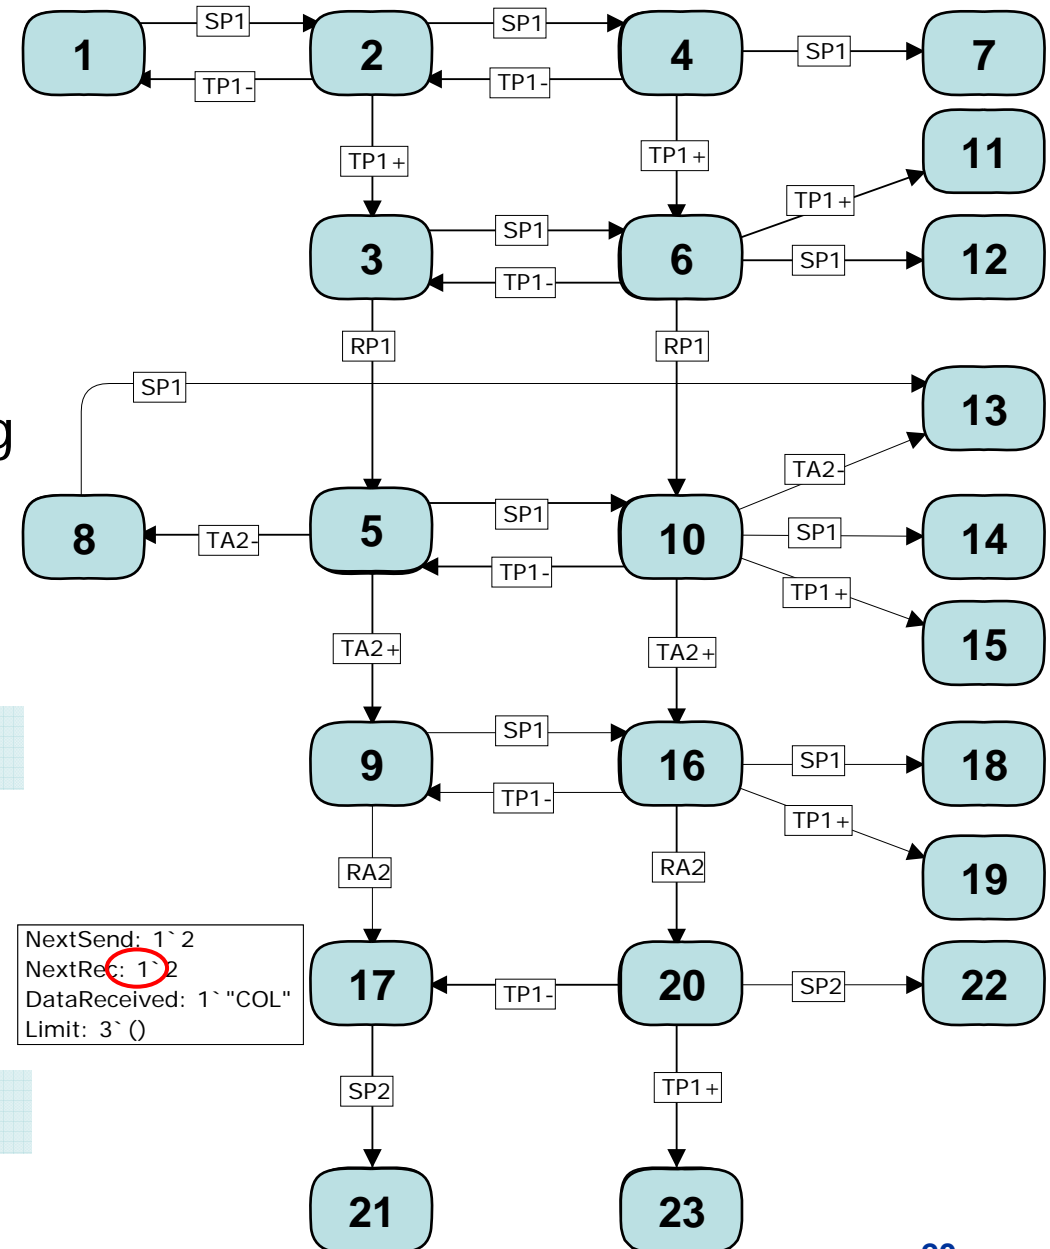
Reachability properties

- The **standard query function** below checks whether marking M_{17} is **reachable** from M_1 – i.e. whether there is a **path** from node 1 to node 17.

Reachable (1,17); true

- We can also check whether M_1 is **reachable** from M_{17} :

Reachable (17,1); false



Reachability properties (SCC)

- It is also possible (and more efficient) to check reachability from the **SCC graph**.
- Then we check whether there exists a **path** from the **SCC** containing the first marking to the **SCC** containing the second marking.

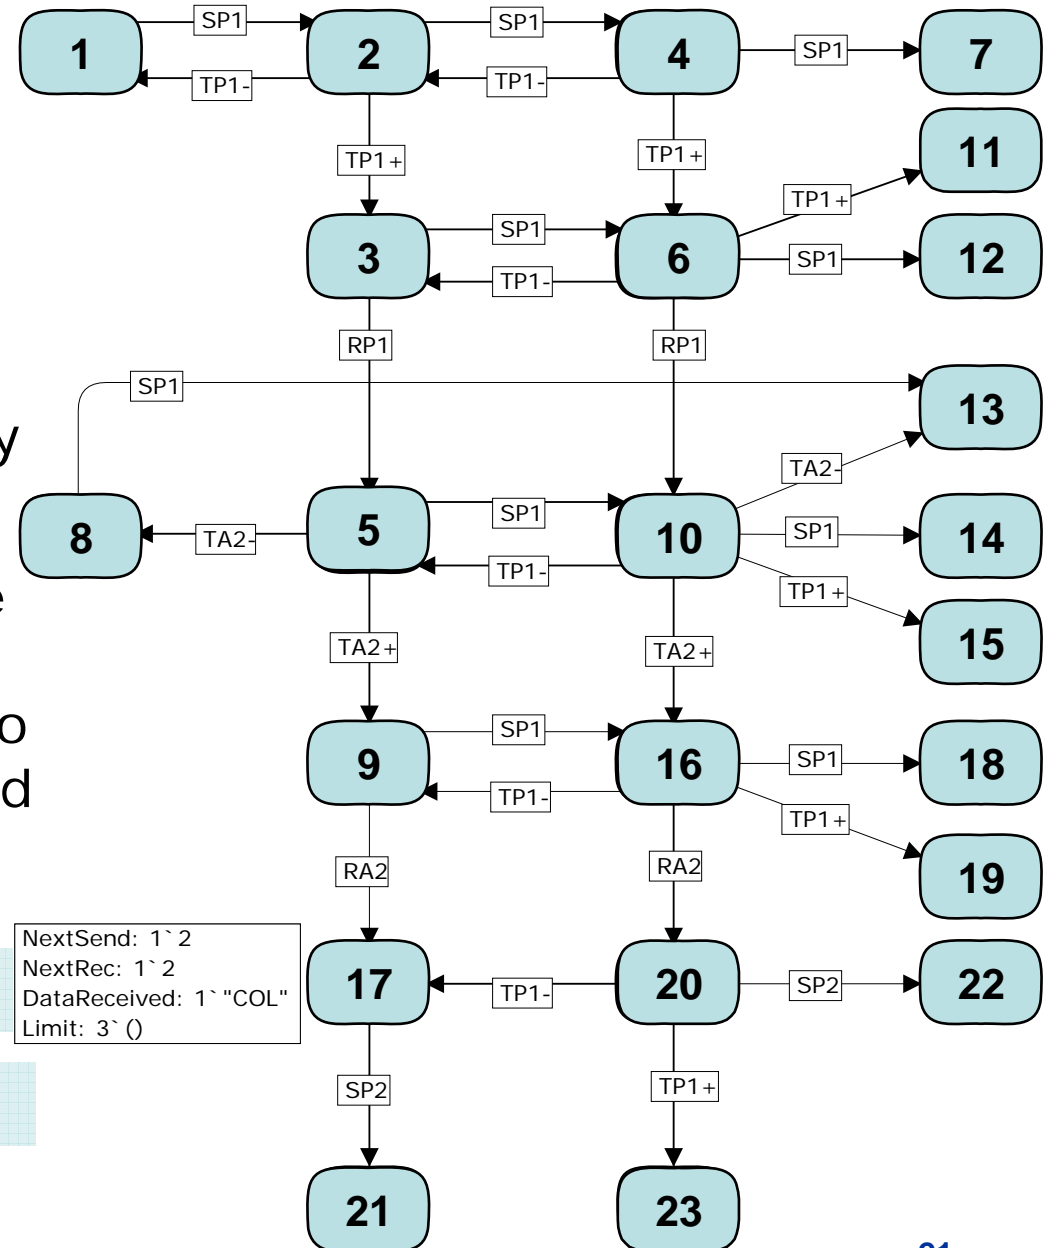
`SccReachable (1,17);`

`true`

`SccReachable (17,1);`

`false`

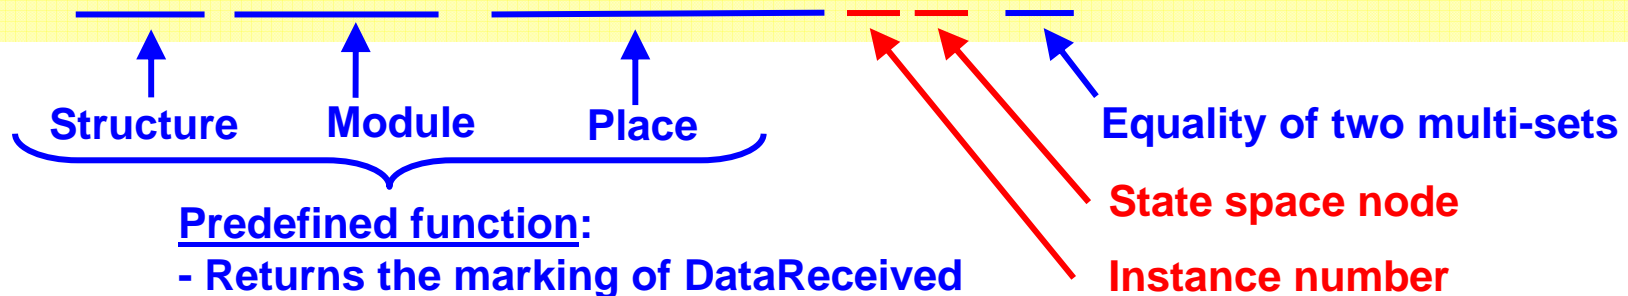
NextSend: 1`2
NextRec: 1`2
DataReceived: 1`"COL"
Limit: 3`()



Desired terminal marking

- The following **predicate** checks whether **node n** represents a marking in which **all data packets** have been **successfully** received.

```
fun DesiredTerminal n =  
  ((Mark.Protocol'PacketsToSend 1 n) == AllPackets) andalso  
  ((Mark.Protocol'NextSend 1 n) == 1'7) andalso  
  ((Mark.Protocol'A 1 n) == empty) andalso  
  ((Mark.Protocol'B 1 n) == empty) andalso  
  ((Mark.Protocol'NextRec 1 n) == 1'7) andalso  
  ((Mark.Protocol'C 1 n) == empty) andalso  
  ((Mark.Protocol'D 1 n) == empty) andalso  
  ((Mark.Protocol'DataReceived 1 n) == 1'"COLOURED PETRI NET")
```



Reachability of desired terminal marking

- The following **query** checks whether the **desired terminal marking** is reachable:

```
ReachablePred DesiredTerminal; true
```

Standard query function:

- Searches through all nodes
- Determines whether some of these fulfil the **predicate**

- It is also possible to find **the node(s)** which represents the **desired terminal marking**:

```
PredAllNodes DesiredTerminal; [4868]
```

Standard query function:

- Searches through all nodes
- Returns a list with those that fulfil the **predicate**



State space report: reachability properties

- The **state space report** does **not** contain information about **reachability properties**.
- The **specific markings** which it is of interest to investigate is highly **model dependent** – and there are **too many** to investigate all pairs.
- The **statistics** in the **state space report** for the protocol shows that there are **more than one SCC**.
- This implies that **not all nodes** in the state space are **mutually reachable** – as demonstrated above using standard query functions.



Integer bounds

- **Integer bounds** counts the **number** of tokens on a place.
- The **best upper integer bound** for a place is the **maximal number** of tokens on the place in a reachable marking.
- The **best lower integer bound** for a place is the **minimal number** of tokens on the place in a reachable marking.
- Places with an **upper integer bound** are **bounded**.
- Places with **no upper integer bound** are **unbounded**.
- 0 is always a **lower integer bound**, but it may **not** be the best.



State space report: integer bounds

Best Integers Bounds	Upper	Lower
PacketsToSend	6	6
DataReceived	1	1
NextSend, NextRec	1	1
A, B, C, D	3	0
Limit	3	0

- **PacketsToSend** has **exactly 6 tokens** in all reachable markings.
- **DataReceived**, **NextSend** and **NextRec** have **exactly one token** each in all reachable markings.
- The **remaining five places** have **between 0 and 3 tokens** each in all reachable markings.



More general integer bounds

- It is also possible to find **integer bounds** for a **set** of places.
- As an example, we might investigate how many tokens we have **simultaneously** on places A and B.

```
fun SumMarkings n =  
  (Mark.StateSpaceProtocol'A 1 n) ++  
  (Mark.StateSpaceProtocol'B 1 n);
```

UpperInteger SumMarkings;

3

LowerInteger SumMarkings;

0

↑
**Standard query
functions**

↑
**Argument must be a function mapping from a
state space node into a multi-set type 'a ms**



More general integer bounds

- It is also possible to investigate **integer bounds** which consider only **certain token colours and places**.
- As an example, we will investigate the minimal and maximal number of tokens with the colour **(1,"COL")** that can **simultaneously** reside on the **places A and B**:

Standard list function:

- Takes a predicate and a list as arguments
- Returns those elements that fulfil the predicate

```
fun SumFirstDataPacket n =  
  (List.filter  
    (fn p => p = (1, "COL"))  
    (SumMarkings n));
```

Marking of places A and B

CPN tools represents multi-sets as lists

```
UpperInteger SumFirstDataPacket;  
LowerInteger SumFirstDataPacket;
```

3

0



Multi-set bounds

- **Integer bounds** count the **number** of tokens **ignoring** the token colours.
- **Multi-set bounds** provide information about the **possible** token colours.
- The **best upper multi-set bound** for a place is a **multi-set** over the colour set of the place.
- The **coefficient** for a **colour c** is the **maximal number** of occurrences of tokens with **colour c** in a reachable marking.
- The **best lower multi-set bound** for a place is a **multi-set** over the **colour set** of the place.
- The **coefficient** for a **colour c** is the **minimal number** of occurrences of tokens with **colour c** in a reachable marking.



State space report: upper multi-set bounds

Best Upper Multi-set Bounds

PacketsToSend	$1'(1, \text{"COL"})++1'(2, \text{"OUR"})++1'(3, \text{"ED "})++1'(4, \text{"PET"})++1'(5, \text{"RI "})++1'(6, \text{"NET"})$
DataReceived	$1''''++1''\text{"COL"}++1''\text{"COLOUR"}++1''\text{"COLOURED " }++1''\text{"COLOURED PET"}++1''\text{"COLOURED PETRI " }++1''\text{"COLOURED PETRI NET"}$
NextSend, NextRec	$1'1++1'2++1'3++1'4++1'5++1'6++1'7$
A, B	$3'(1, \text{"COL"})++3'(2, \text{"OUR"})++3'(3, \text{"ED "})++3'(4, \text{"PET"})++3'(5, \text{"RI "})++3'(6, \text{"NET"})$
C, D	$3'2++3'3++3'4++3'5++3'6++3'7$
Limit	$3'()$

- The **upper bound** for **DataReceived** is a multi-set with **seven elements** although the place always has **exactly one token**.



State space report: lower multi-set bounds

Best Lower Multi-set Bounds

PacketsToSend	$1'(1, \text{"COL"})++1'(2, \text{"OUR"})++1'(3, \text{"ED "})++1'(4, \text{"PET"})++1'(5, \text{"RI "})++1'(6, \text{"NET"})$
DataReceived	empty
NextSend, NextRec	empty
A, B, C, D	empty
Limit	empty

- The **lower bound** for **DataReceived** is **empty** although the place always has **exactly one token**.



More general multi-set bounds

- Upper and lower multi-set bounds can be **generalised** to **sets of places** in a similar way as described for integer bounds.

```
UpperMultiSet SumMarkings;  
LowerMultiSet SumMarkings;
```

↑
Standard query functions

↑
Argument must be a function mapping from a state space node into a multi-set type 'a ms

```
3 `(1, "COL") ++ 3 `(2, "OUR") ++ 3 `(3, "ED ") ++  
3 `(4, "PET") ++ 3 `(5, "RI ") ++ 3 `(6, "NET")
```

```
empty
```



More general multi-set bounds

- Upper and lower multi-set bounds can also be **generalised** to **specific token colours** residing on a set of places in a similar way as described for integer bounds.

UpperMultiSet SumFirstDataPacket;

LowerMultiSet SumFirstDataPacket;

3 ` (1, "COL")

empty

↑
**Standard query
functions**

↑
**Argument must be a function mapping from a
state space node into a multi-set type 'a ms**



Integer and multi-set bounds

- The two kinds of bounds **supplement** each other and provides **different** kinds of information.

DataReceived	1
--------------	---

Tells us that DataReceived has **at most one token**, but gives us no information about the token colours.


DataReceived	1""++1"COL"++1"COLOUR"++1"COLOURED "++ 1"COLOURED PET"++1"COLOURED PETRI "++ 1"COLOURED PETRI NET"
--------------	--

Tells us that DataReceived can have **seven different token colours**, but not whether they can be present simultaneously.



Home marking

- A **home marking** is a marking M_{home} which can be reached from **any** reachable marking.



- This means that it is **impossible** to have an occurrence sequence which **cannot be extended** to reach M_{home} .
- The **home property** tells that it is **possible** to reach M_{home} .
- However, there is **no guarantee** that this will happen.

State space report: home markings

Home Properties

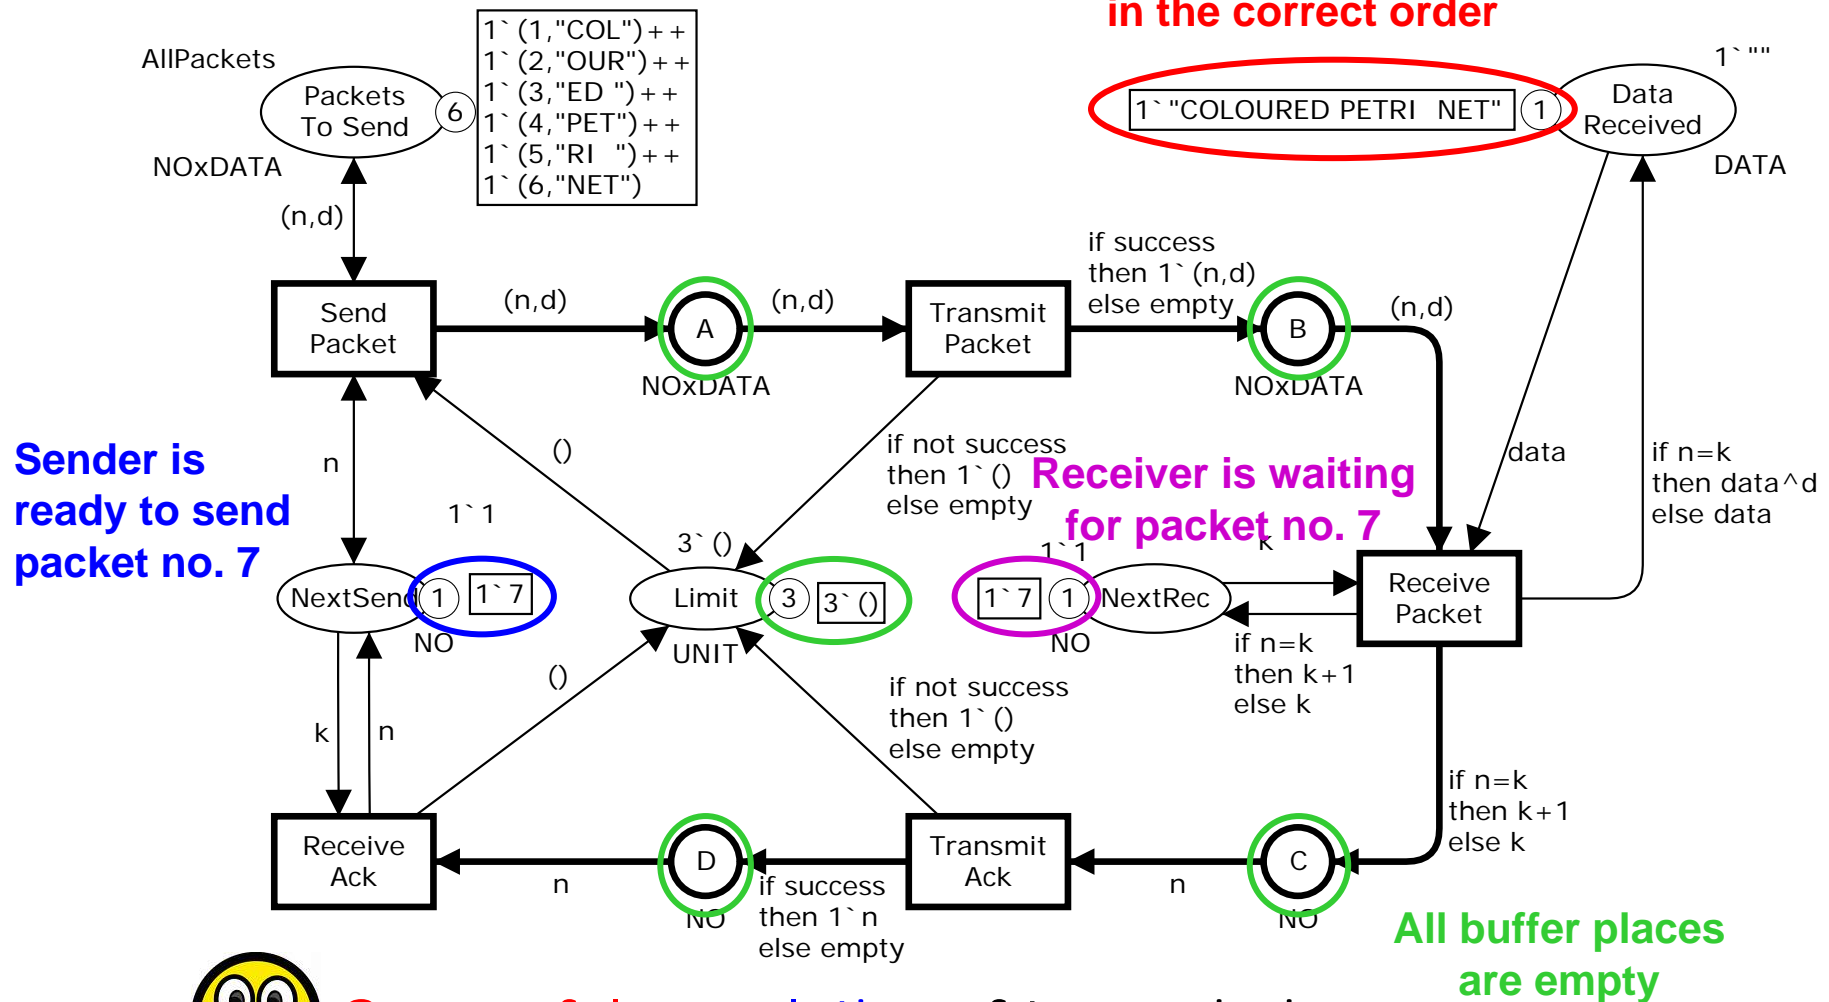
Home Markings: [4868]

- There is a **single home marking** represented by node number 4868.
- The **marking** of this node can be shown in the **CPN simulator**.



Home marking

All packets have been received
in the correct order




Successful completion of transmission.



Home space


- A **home space** is a **set** of markings M^*_{home} such that at least one marking in M^*_{home} can be reached from **any** reachable marking.



- This means that it is **impossible** to have an occurrence sequence which **cannot be extended** to reach a marking in M^*_{home} .
- The **home property** tells that it is **possible** to reach a marking in M^*_{home} .
- However, there is **no guarantee** that this will happen.

Home predicate

- A **home predicate** is a **predicate** on markings $\text{Pred}_{\text{home}}$ such that at least one marking satisfying $\text{Pred}_{\text{home}}$ can be reached from **any** reachable marking.



- This means that it is **impossible** to have an occurrence sequence which **cannot be extended** to reach a marking satisfying $\text{Pred}_{\text{home}}$.
- The **home property** tells that it is **possible** to reach a marking satisfying $\text{Pred}_{\text{home}}$.
- However, there is **no guarantee** that this will happen.

Use of home predicate

- Instead of inspecting node 4868 in the CPN simulator we can check whether `DesiredTerminal` is a **home predicate**:

```
fun DesiredTerminal n =  
  ((Mark.Protocol'PacketsToSend 1 n) == AllPackets) andalso  
  ((Mark.Protocol'NextSend 1 n) == 1'7) andalso  
  ((Mark.Protocol'A 1 n) == empty) andalso  
  ((Mark.Protocol'B 1 n) == empty) andalso  
  ((Mark.Protocol'NextRec 1 n) == 1'7) andalso  
  ((Mark.Protocol'C 1 n) == empty) andalso  
  ((Mark.Protocol'D 1 n) == empty) andalso  
  ((Mark.Protocol'Data_Received 1 n) == 1'"COLOURED PETRI NET")
```

HomePredicate DesiredTerminal;

true

↑
Standard query function

↑
Argument must be a predicate on markings





Use of home properties to locate errors

- Home properties are excellent to locate certain kinds of errors.
- As an example, consider a CPN model of a telephone system.
- If all users stop calling and terminate all ongoing calls, the system is expected to reach an idle system state in which all lines and all equipment are unused and no calls are in progress.
- The idle system state will be represented:
 - by a home marking (if the system is without memory)
 - by a home space / home predicate (if information is stored about prior activities).
- If one or more reachable markings exist from which we cannot reach the idle system state, we may have made a modelling error or a design error – e.g., forgotten to return some resources.



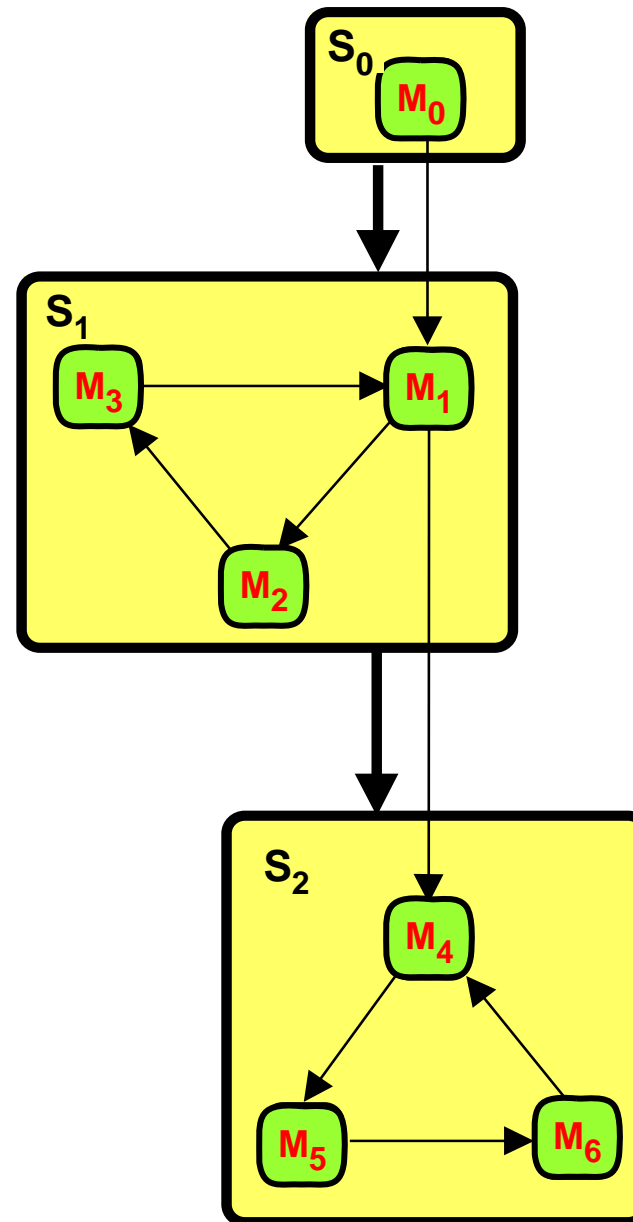
Home markings and SCCs

- The existence of **home markings** can be **determined** from the **number** of **terminal SCCs**.
- **Only one** terminal SCC:
 - **All markings** in the **terminal SCC** are home markings.
 - **No other markings** are home markings.
- **More than one** terminal SCC:
 - **No** home markings.



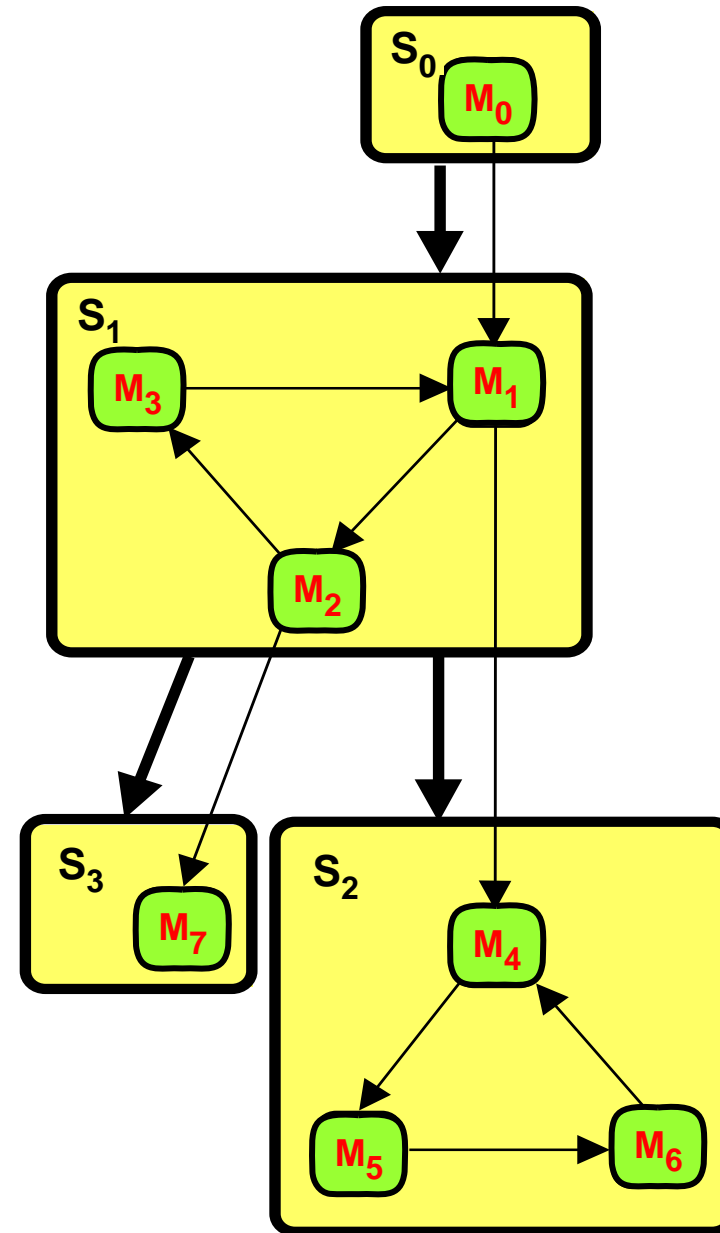
Single terminal SCC

- All markings in the terminal SCC S_2 are home markings.
- No other markings are home markings.




More than one terminal SCC

- No home markings.
- When one of the terminal SCCs S_2 and S_3 has been reached, it is impossible to leave it again.



Single SCC

- All reachable markings are home markings.
- They are mutually reachable from each other.



Calculation of home markings

- The **CPN state space tool** uses the following **query** to calculate the set of all **home markings**:

```
fun ListHomeMarkings () =  
  let  
    val Terminal_Sccs = PredAllSccs SccTerminal;  
  in  
    case Terminal_Sccs of  
      [scc] => SccToNodes scc  
    | _ => []  
  end;
```

Exactly one terminal SCC →

Checks whether an SCC is terminal →


Standard query function:
- Searches through all nodes in the SCC graph
- Returns those which fulfil the **predicate**

Returns the state space nodes in the strongly connected component scc →



Home spaces and SCCs

- The size of **home spaces** can be **determined** from the **number** of **terminal** components in the **SCC graph**.
- A set of markings is a **home space** if and only if it **contains a node** from **each terminal SCC**.
- Home spaces must have **at least as many elements** as there are **terminal SCCs**.
- Each **home marking** is a **home space** with only one element.
- A system may have **home spaces without** having **home markings**.



Liveness properties – being dead

- A marking M is **dead** if M has no enabled transitions.
- A transition t is **dead** if t never can occur – i.e. is disabled in all reachable markings.
- Generalisations:
- A binding element is **dead** if it can never become enabled.
- A set of binding elements is **dead** if none of the binding elements can become enabled.
- A set of transitions is **dead** if the union of their binding elements is dead.



State space report: being dead

Liveness Properties

Dead Markings: [4868]

Dead Transitions: None

Live Transitions: None

- There is a **single dead marking** represented by node number 4868.
 - **Same marking** as home marking.
- There are **no dead transitions**.



Marking no 4868

- We have seen that marking M_{4868} represents the state in which we have achieved **successful completion** of the transmission.
- M_{4868} is the **only dead marking**.
- Tells us that the system is **partially correct**. If execution terminates we will have the correct result.
- M_{4868} is a **home marking**.
- Tells us that it **always** is **possible** to reach the **correct result** – independently of the number of losses and overtakings.



Being dead

- It is straightforward to check whether markings, transitions and binding elements are dead.
- A marking is dead if the corresponding state space node has no outgoing arcs.
- A transition is dead if it does not appear on an arc in the state space.
- A binding element is dead if it does not appear on an arc in the state space.
- A set of binding elements is dead if no binding element in the set appears on an arc in the state space.
- A set of transitions is dead if none of their binding elements appear on an arc in the state space.



Calculation of dead markings

- The **CPN state space tool** uses the following **query** to calculate the set of all **dead markings**:

```
fun ListDeadMarkings () =  
  PredAllNodes (fn n => (OutArcs n) = []);
```

Maps a state space node
into its outgoing arcs

Standard query function:
- Searches through all
nodes in the state space
- Returns a list with those
that fulfil the **predicate**

**Checks whether the set of
output arcs is empty**



Calculation of dead transitions

- The **CPN state space tool** uses the following **query** to check whether a **transition instance** is **dead**:

```
fun TransitionInstanceDead ti =  
  (PredAllArcs (fn a => ArcToTI a = ti)) = [];
```

Maps a state space arc into
its transition instance

Standard query function:

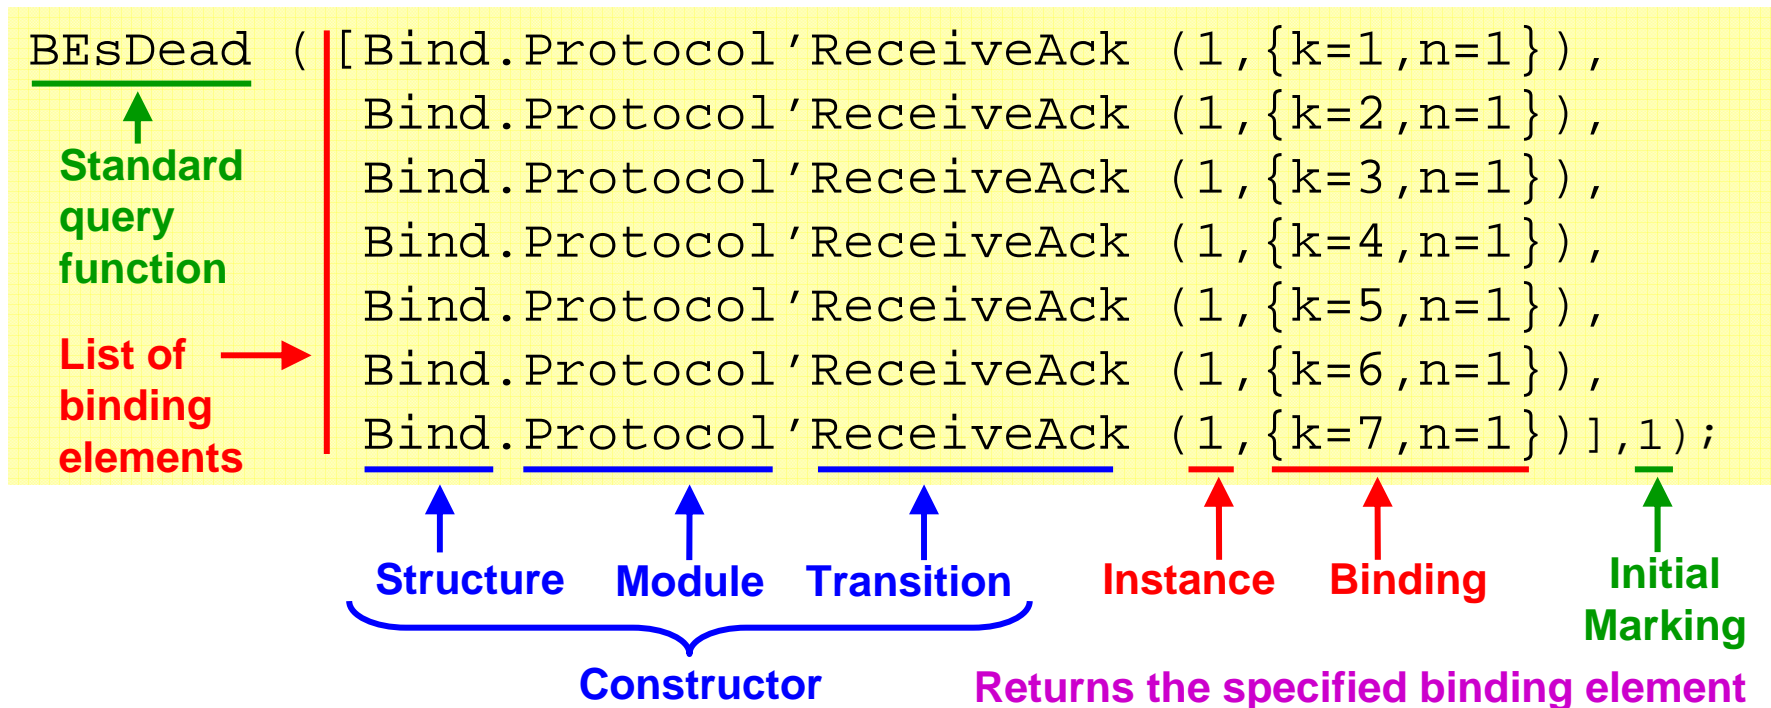
- Searches through all arcs in the state space
- Returns a list with those that fulfil the **predicate**

Checks whether the arc a has the
transition instance ti in its label



Calculation of dead binding elements

- We want to **check** whether the **Sender** can receive an **acknowledgement** with **sequence number 1**.



true

Not possible to receive such acknowledgments.




Liveness properties – being live

- A transition t is **live** if we from all reachable markings can find an **occurrence sequence** containing t .



- **Liveness** tells that it is **possible** for t to occur.
- However, there is **no guarantee** that this will happen.

Liveness is a strong property



- If the live transition t occurs in the marking M_2 we reach another reachable marking.
- We can use the new marking as M_1 and hence t is able to occur once more, and so on.
- This means that there exists infinite occurrence sequences from M_1 in which t occurs infinitely many times.
- It is possible to be non-dead without being live.

State space report: being live

Liveness Properties

Dead Markings: [4868]

Dead Transitions: None

Live Transitions: None

- There are **no live transitions**
- **Trivial consequence** of the existence of a **dead marking**.




Generalisations of liveness

- A **binding element** is **live** if it can always become enabled.
- A **set of binding elements** is **live** if it is always possible to enable at least one binding element in the set.
- A **set of transitions** is **live** if the union of their binding elements is live.



Liveness properties and SCCs

- **Liveness** can be determined from the **SCC graph**.




- A **transition/binding element** is **live** if and only if it appears on at least one arc in each terminal SCC.
- A **set of transitions/binding elements** is **live** if and only if each of the terminal SCCs contains at least one arc with a transition/binding element from the set.




Single terminal SCC

- A transition is **live** if it appears on an **arc** in the **terminal SCC** S_2 .




More than one terminal SCC

- No live transitions.
- S_3 is terminal and trivial.
- M_7 is a dead marking.




Single SCC

- A transition is **live** if it appears on an **arc** in the SCC.
- A transition is **live** if and only if it is **non-dead**.



Fairness properties

- A transition t is **impartial** if t occurs **infinitely** often in **all** infinite occurrence sequences.



State space report: fairness properties

Fairness Properties

Impartial Transitions: [SendPacket 1, TransmitPacket 1]

Instance

Instance

- SendPacket and TransmitPacket are impartial.
- If one of these are removed (or blocked by the guard false) the protocol will have no infinite occurrence sequences.
- The other three transitions are not impartial.
- If we remove the Limit place only SendPacket will be impartial.
- Adding the Limit place has changed the behavioural properties.



Generalisations of impartial

- A **binding element** is **impartial** if it occurs infinitely often in all infinite occurrence sequences.
- A **set of binding elements** is **impartial** if binding elements from the set occurs infinitely often in all infinite occurrence sequences.
- A **set of transitions** is **impartial** if the union of their binding elements is impartial.



Fairness properties and SCCs

- Impartiality of a transition/binding element can be checked by means of an SCC graph:
 1. Construct the pruned state space in which all arcs with appearances of the transition/binding element are removed.
 2. Construct the SCC graph of the pruned state space.
 3. Check whether the SCC graph for the pruned state space has the same number of nodes and arcs as the state space for the pruned state space.
 4. If this is the case the pruned state space has no cycles.
 5. This implies that all cycles in the original state space contain an arc with an appearance of the transition/binding element.
 6. Hence the transition/binding element is impartial.
- Impartiality of a set of transitions/binding elements is checked in a similar way.

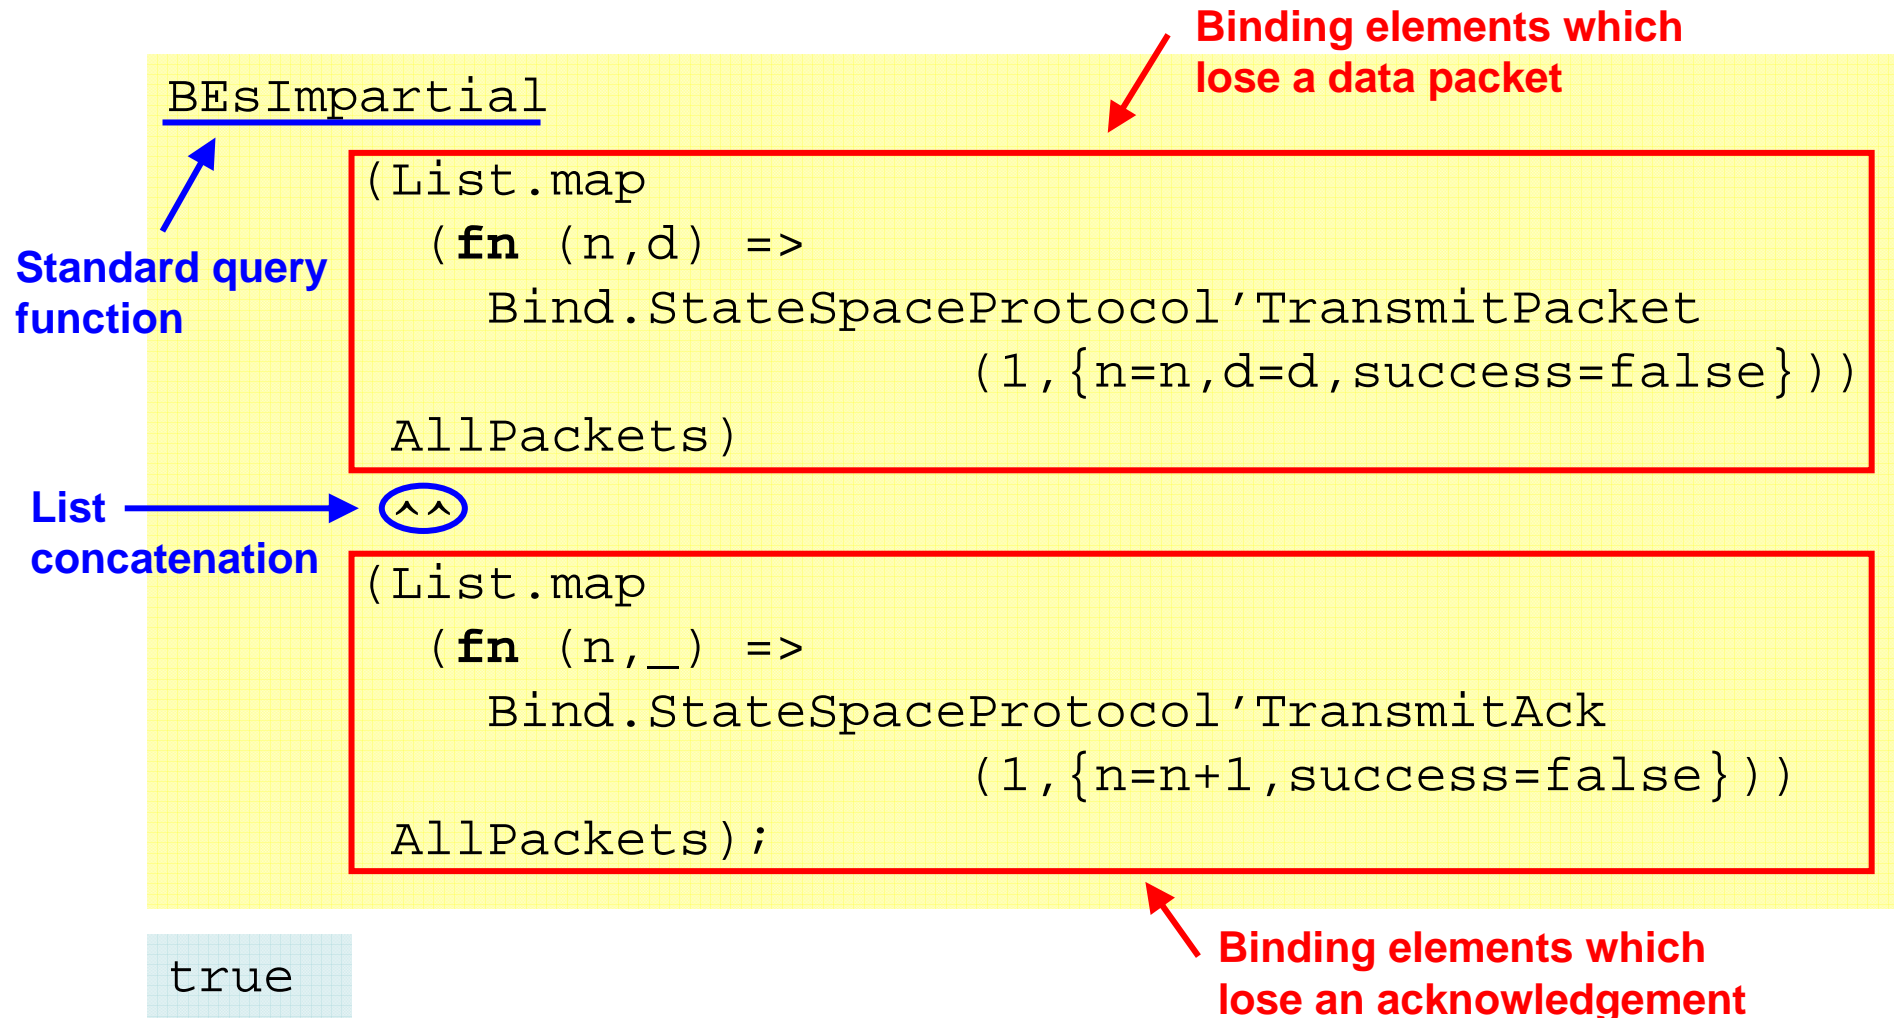


Use of fairness properties

- As an example, we will investigate whether the set of binding elements corresponding to **loss** of **data packets** and **acknowledgements** is impartial.
- If the protocol does **not terminate** we expect this to be because the network keeps **losing packets**, and we therefore expect this set of binding elements to be **impartial**.



Use of fairness properties



Query functions – summary

- The **state space report** contains information about **standard behavioural properties** which make sense for **all** CPN models.
- **Non-standard behavioural properties** can be investigated by means of **queries**.
- For some purposes it is sufficient to provide **arguments** to a **predefined query function** – e.g. to check whether a set of markings constitute a home space.
- For other more special purposes it is necessary to write your own **query functions** using the **CPN ML programming language**.



Example of user-defined query function

- We want to **check** whether the protocol obeys the **stop-and-wait** strategy – i.e. that the sender always sends the data packet **expected** by the receiver (or the previous one)

```
fun StopWait n =  
  let  
    val NextSend = ms_to_col (Mark.Protocol'NextSend 1 n);  
    val NextRec   = ms_to_col (Mark.Protocol'NextRec 1 n);  
  in  
    (NextSend = NextRec) orelse (NextSend = NextRec - 1)  
  end;
```

Converts a multi-set 1`x with
one element to the colour x

```
val SWviolate = PredAllNodes (fn n => not(StopWait));
```

Negation

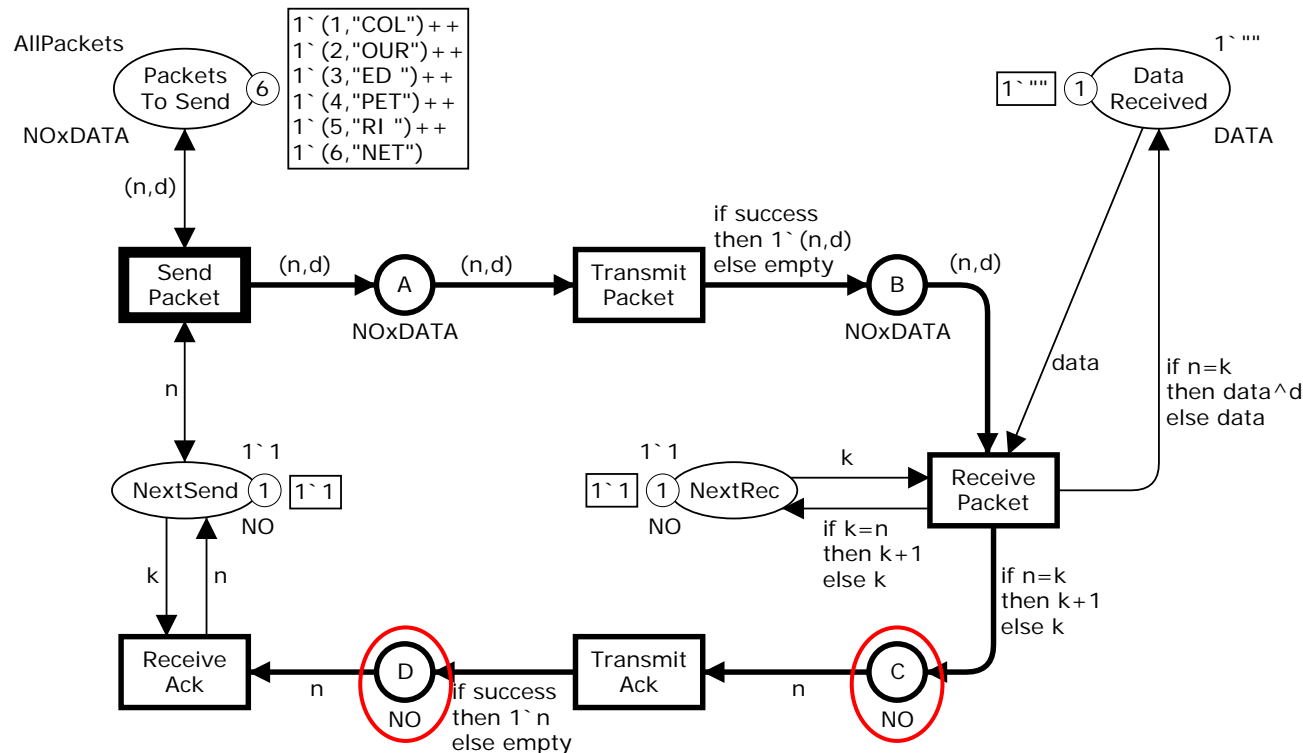
- The **stop-and-wait** strategy is **not** satisfied (7020 violations).

We check whether some states
violate the property.
This is easier than checking
that all states fulfil the property.



Violation of stop-and-wait strategy

- Acknowledgements may **overtake** each other on C and D.
- This means that it is possible for the sender to receive an **old acknowledgement** which **decrements** NextSend.

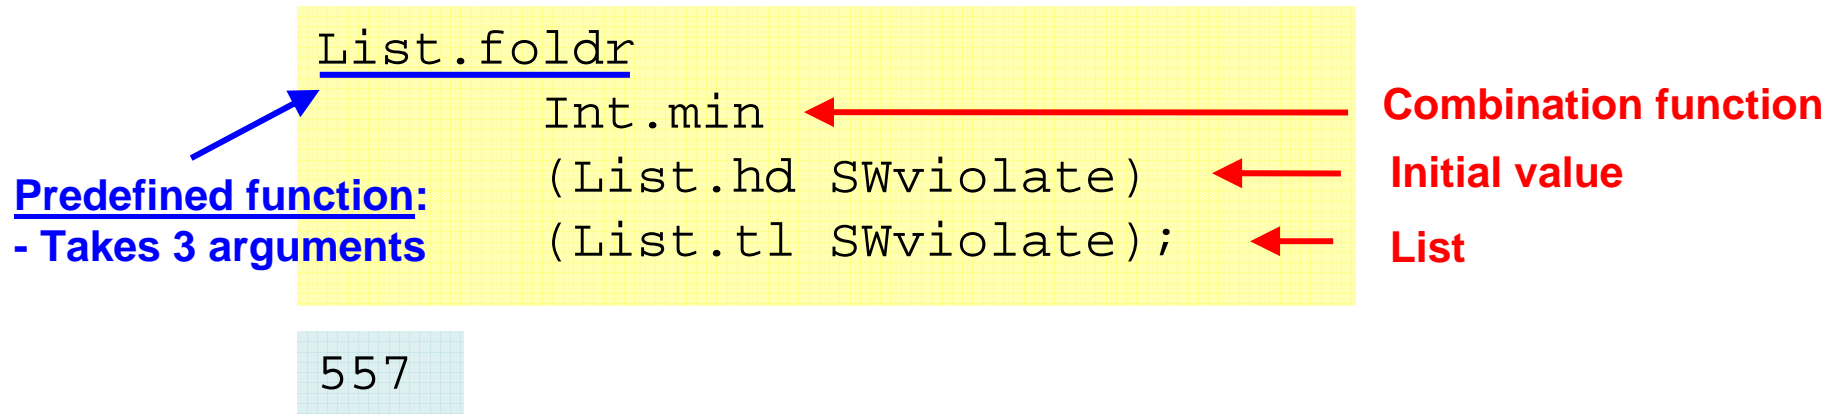


Shortest counter example

- We want to construct a **shortest counter example** – i.e. to find one of the **shortest occurrence sequences** leading from the initial marking to a marking where the **predicate does not hold**.
- The state space is generated in **breadth-first order**.
- Hence, we search for the **lowest numbered** node in the list **SWviolate**.



Lowest node in SWviolate



- The function **iterates** over the list.
- In each iteration the **combination function** is applied to the pair consisting of the **current** element in the list and the value returned by the **previous** application of the **combination function**.
- In the **first iteration**, the **initial value** plays the role of the result from the previous application.

Shortest counter example

```
ArcsInPath(1, 557);
```

Predefined function:

- Returns the arcs in one of the shortest paths from 1 to 557

```
[1, 3, 9, 16, 27, 46, 71, 104,  
142, 201, 265, 362, 489, 652,  
854, 1085, 1354, 1648]
```

↑
18 arcs

- The **path** can be **visualised** using the **drawing facilities** in the CPN state space tool.
- This is the **same drawing facilities** that were used to **visualise** the initial fragment of the state space (at the beginning of this lecture).



Bindings elements in counter example

- The **binding elements** in the **shortest path** can be obtained by the following query:

```
List.map (ArcToBE (ArcsInPath(1,557))) ;
```

**Maps a state space arc
into its binding element**

**Shortest path with
counter example**



Shortest counter example

Packet no 1 and its ack	●	1	(SendPacket, <d="COL",n=1>
	●	2	(TransmitPacket, <n=1,d="COL",success=true>
	●	3	(ReceivePacket, <k=1,data="",n=1,d="COL">
	●	4	(SendPacket, <d="COL",n=1>
	●	5	(TransmitAck, <n=2,success=true>
	●	6	(ReceiveAck, <k=1,n=2>
Packet no 2 and its ack	●	7	(SendPacket, <d="OUR",n=2>
	●	8	(TransmitPacket, <n=1,d="COL",success=true>
	●	9	(TransmitPacket, <n=2,d="OUR",success=true>
	●	10	(ReceivePacket, <k=2,data="COL",n=1,d="COL">
	●	11	(ReceivePacket, <k=2,data="COL",n=2,d="OUR">
	●	12	(TransmitAck, <n=3,success=true>
Packet no 3 NextRec = 4	●	13	(ReceiveAck, <k=2,n=3>
	●	14	(SendPacket, <d="ED ",n=3>
	●	15	(TransmitPacket, <n=3,d="ED ",success=true>
	●	16	(ReceivePacket, <k=3,data="COLOUR",n=3,d="ED ">
Retrans- mission NextSend = 2	●	17	(TransmitAck, <n=2,success=true>
	●	18	(ReceiveAck, <k=3,n=2>



State space for revised protocol

- The **state space** contains **1,823 nodes** and **6,829 arcs**.
- **Before** we had **13,215 nodes** and **52,874 arcs**.
- As before there is a **single dead marking** which corresponds to the **desired** terminal marking, where all packets have been successfully transmitted.
- The **new protocol** is **partially correct**.
- Now there are **no home markings**.
- We can reach situations from which it is **impossible** to reach the **desired** terminal marking.
- The **new protocol** may enter a **live-lock**.



Analysis of revised protocol

- The **dead marking** is no longer a **home marking** and hence we must have one or more terminal SCCs from which we **cannot** reach the dead marking.
- These **terminal SCCs** can be found by the following **query** which returns all SCCs that are **terminal** but **not trivial**:

```
PredAllSccs (fn scc => SccTerminal scc andalso  
              not (SccTrivial scc));
```

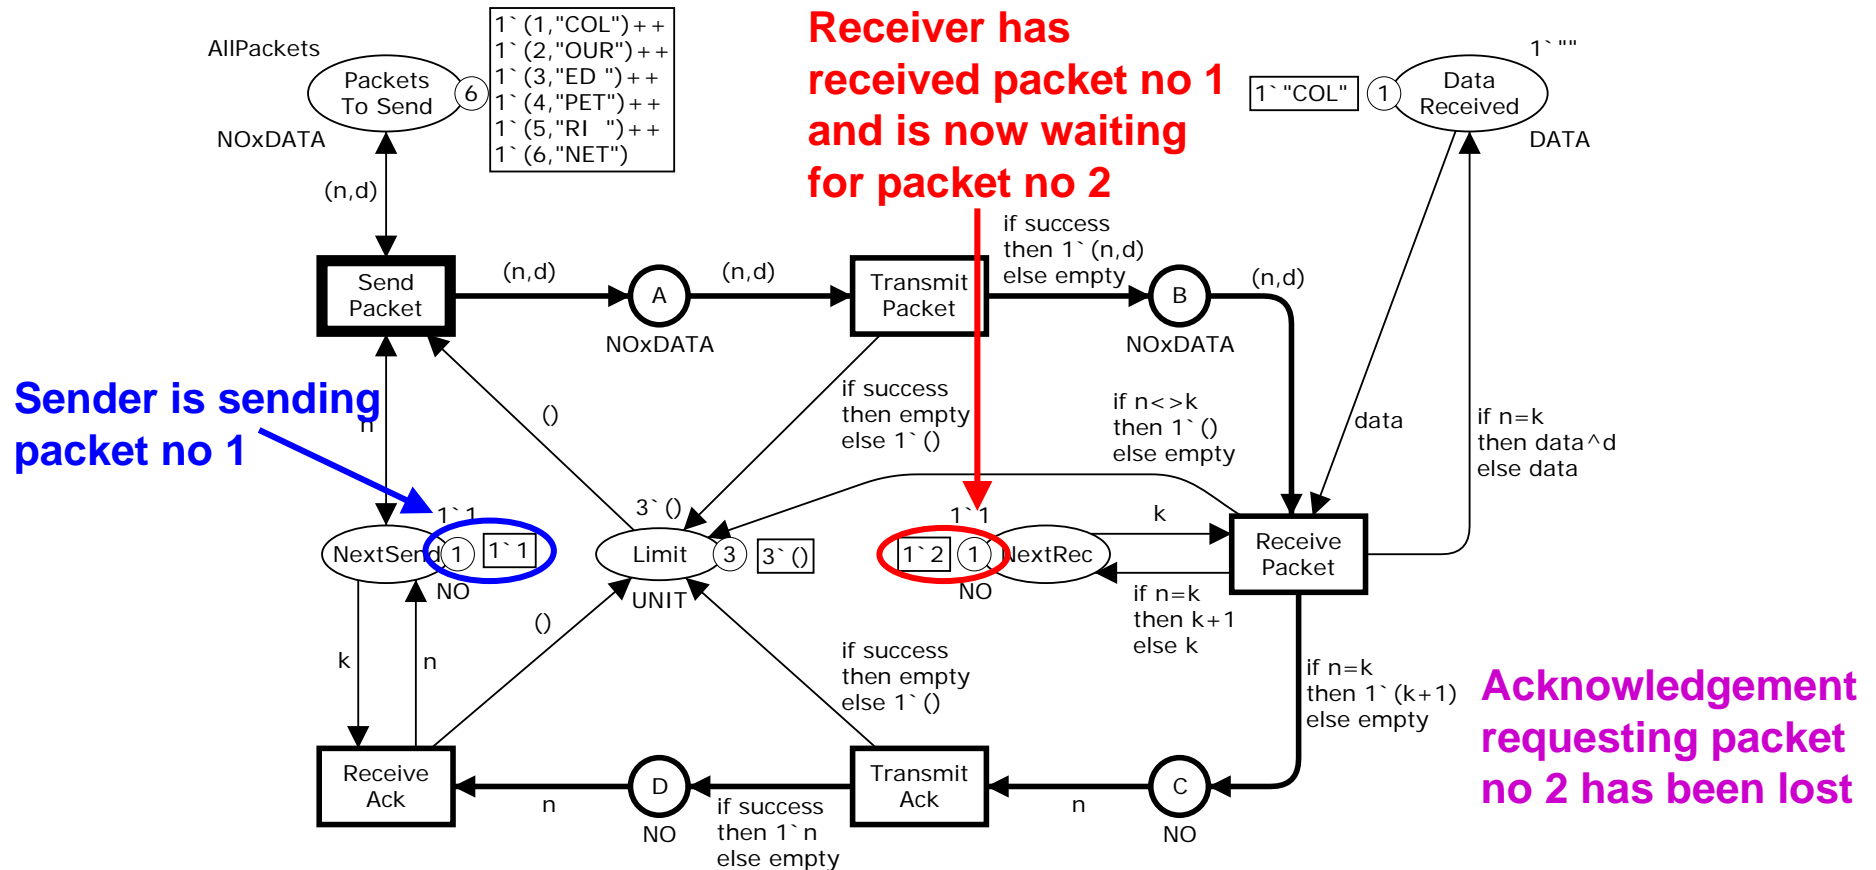
↖ **Standard query function:**

- Searches through all nodes in the SCC graph
- Returns those which fulfil the **predicate**

- The result of the **query** is a list with **six SCCs**.
- The state space **nodes** in the six SCCs can be obtained using the function **SccToNodes**.
- To get a shortest counter example, we choose the **lowest numbered** node which is **node 12**.



Marking no 12



All data packets will be "wrong".
No acknowledgements will be sent – no progress.

What is wrong

- The **analysis** of **marking no 12** has told us what the problem is.
- The sender continues to send **wrong packets** and the receiver never sends an acknowledgement which can **correct** the problem.
- We might also want to know **how we arrived** at this unfortunate situation.
- This is done by constructing an **error trace** / **counter example**.



Counter example

- The **query** below returns a **list** with all the **arcs** in one of the **shortest paths** from node 1 (initial marking) to node number 12:

```
ArcsInPath(1,12);
```

- The **binding elements** in the **shortest path** can be obtained by the following query:

```
List.map (ArcToBE (ArcsInPath(1,12))) ;
```

Maps a state space arc
into its binding element



Counter example

- The result of the **query** is the following **list of binding elements**:

```
1  (SendPacket, <d="COL", n=1>)  
2  (TransmitPacket, <d="COL", n=1, success=true>)  
3  (ReceivePacket, <d="COL", n=1, k=1, data="">)  
4  (TransmitAck, <n=2, success=false>)
```

- We see that **data packet no 1** was sent, successfully transmitted, and received.
- However, the **acknowledgment** requesting **data packet no 2** was **lost** on the network.



System configurations

- With **state space analysis** we always investigate a system for a **particular configuration** of the **system parameters**.
- **In practice** it is often sufficient to consider a **few rather small configurations** – although we **cannot** be totally sure that larger configurations will have the same properties.
- As **system parameters** increase the **size of the state space** increases – often in an **exponential way**.
- This is called the **state space explosion**, and it is one of the most **severe limitations** of the state space method.



Different system configurations

Limit	Packets	Nodes	Arcs	Limit	Packets	Nodes	Arcs
1	1	9	11	5	1	217	760
1	2	17	22	5	3	17,952	97,963
1	200	1,601	2,200	5	5	269,680	1,655,021
1	600	4,801	6,600	7	1	576	2,338
2	1	26	53	7	2	11,280	64,297
2	5	716	1,917	7	3	148,690	1,015,188
2	50	93,371	258,822	10	1	1,782	8,195
2	140	746,456	2,072,682	10	2	76,571	523,105
3	1	60	159	12	1	3,276	15,873
3	5	7,156	28,201	12	2	221,117	1,636,921
3	10	70,131	286,746	13	1	4,305	21,294
3	20	622,481	2,583,361	13	2	357,957	2,737,878



Is it worthwhile?

- State space analysis can be a **time consuming** process where it takes **many hours** to generate the state spaces and verify the desired properties.
- However, it is **fully automatic** and hence requires much **less human work** than lengthy simulations and tests.
- It may **take days** to verify the properties of a system by means of state spaces.
- However, this is still a relatively **small investment**:
 - compared to the **total number of resources** used in a system development project.
 - compared to the cost of **implementing, deploying** and **correcting** a system with **errors** that could have been detected in the design phase.



Partial state spaces

- It is sometimes **impossible** to generate the full state space for a given **system configuration** – either because it is too big or takes too long time.
- This means that only a **partial state space** – i.e. a fragment of the state space is generated.
- Partial state spaces **cannot** in general be used to **verify properties**, but they may **identify errors**.
- As an example, an **undesirable dead marking** in a partial state space will **also be present** in the full state space.
- **Partial state spaces** can in that sense be viewed as being positioned **between simulation** and **state spaces**.
- The CPN state space tool has a **number of parameters** to control the generation of **partial state spaces**.



State spaces - summary

- State spaces are **powerful** and **easy** to use.
 - **Construction** and **analysis** can be **automated**.
 - The user do **not** need to know the **mathematics** behind the analysis methods.
- The main drawback is the **state explosion** – i.e. the **size** of the state space.
 - The present CPN state space tool handles state spaces with up to **one million** states.
 - For many systems this is **not sufficient**.
 - A **much more efficient** state space tool is under development.



Reduced state spaces

- Fortunately, it is often possible to construct **reduced** state spaces – **without losing analytic power**.
- This is done by exploiting:
 - **Progress measure**.
 - **Symmetries** in the modelled system.
 - Other kinds of **equivalent** behaviour.
 - **Concurrency** between events.
- The **reduction methods** rely on **complex mathematics**.
- An **overview** of the **advanced state space methods** is given in Chapter 8.

