

Coloured Petri Nets

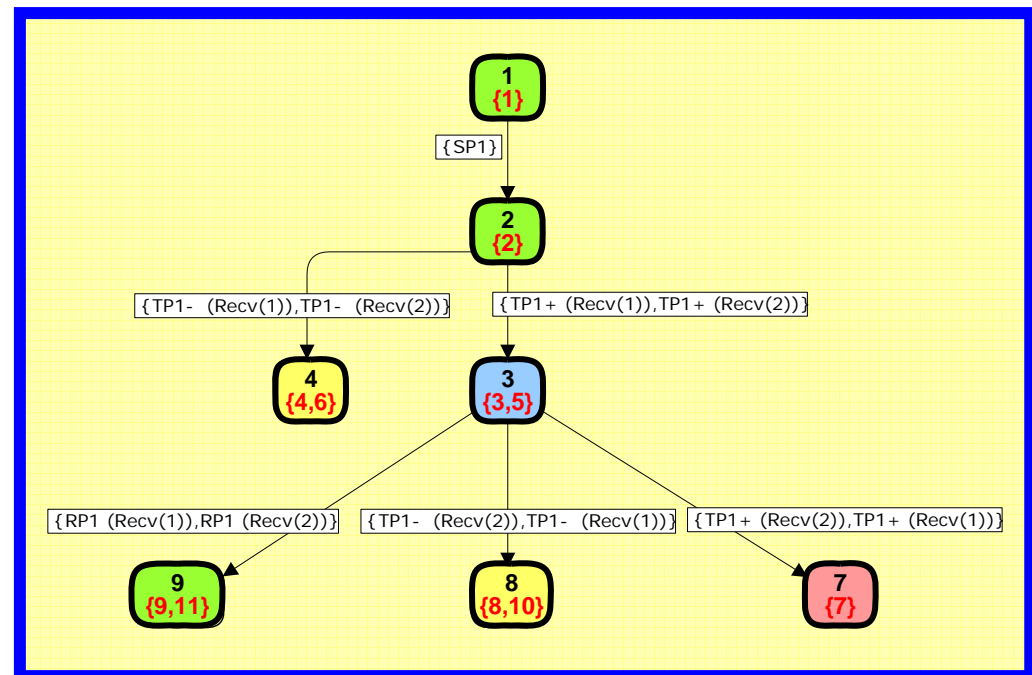
Modelling and Validation of Concurrent Systems

Chapter 8: Advanced State Space Methods

Kurt Jensen &
Lars Michael Kristensen

{kjensen,lmkristensen}
@daimi.au.dk

© February 2008



UNIVERSITY OF AARHUS

Modelling and Validation of Distributed Systems Group
Department of Computer Science

Kurt Jensen and Lars M. Kristensen
Coloured Petri Nets

State space methods

- The main **limitation** of using state spaces to verify behavioural properties of systems is the **state explosion problem**.
- State spaces of many systems have an **astronomical number** of **reachable states**.
- This means that they are **too large** to be handled with the available computing power:
 - **Memory**.
 - **CPU speed**.



State space reduction methods

- Methods for **alleviating** the state explosion problem is an active area of research. They allow:
 - **faster** construction,
 - more **compact** representation (less memory).
- A **large collection** of state space reduction methods exists.
- The **reduction methods** have significantly **increased** the class of systems that can be verified **in practice**.
- State spaces can now be used to verify systems of **industrial size**.



Independent of modelling language

- Most state space reduction methods are **independent** of the **concrete modelling language** and hence applicable for a large class of such languages (e.g. all transition systems).
- Some of the reduction methods have been developed **within** the **context of the CPN modelling language**:
 - Sweep-line method.
 - Symmetry method.
 - Equivalence method.
- Other reduction methods have been developed **outside** the context of the CPN modelling language.



Why different reduction methods?

- State space reduction methods typically exploit certain **characteristics** of the system under analysis.
- **No single reduction method** works well for all kind of systems.
- Furthermore, the methods often **limit** the **verification questions** that can be answered.
- When verifying a **concrete system** one must therefore choose a method that:
 - exploits **characteristics present** in the system.
 - preserves the **behavioural properties** to be verified.



On-the-fly verification

- Many reduction methods are based on the paradigm of **on-the-fly verification**.
- The **verification question** is stated **before** the exploration of the state space starts.
- The **state space exploration** is done **relative** to the provided verification question.
- This makes it possible to **terminate** the state space exploration as soon as the **answer** to the verification question has been obtained – ignoring irrelevant parts.



Model checking

- Many advanced state space reduction methods use **temporal logic** for stating the verification questions :
 - **Linear-time** temporal logic (**LTL**).
 - **Computation tree** temporal logic (**CTL**).
- The use of **temporal logic** for stating and checking verification questions is referred to as **model checking**.



State spaces are kept in main memory

- The amount of available **main memory** is often the **limiting factor** in the practical use of state spaces.
- During **construction** of the state space, the set of markings encountered are kept in **main memory**.
- This allows us to **recognise already visited markings** and thereby ensure that the state space exploration **terminates**.

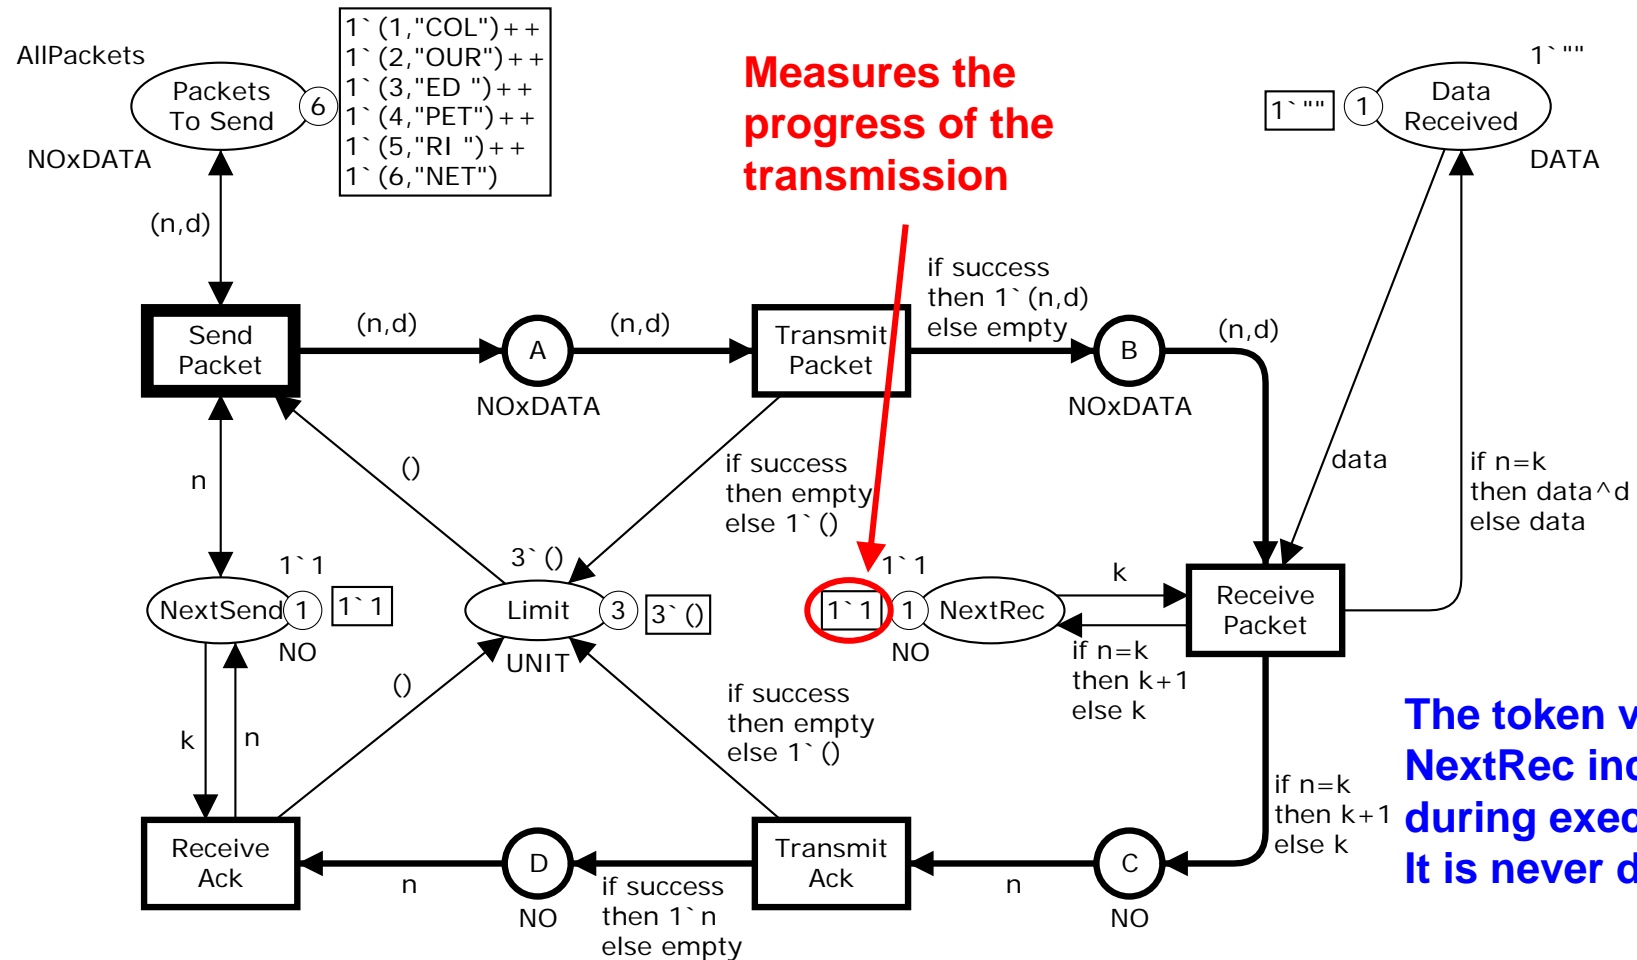


Method 1: Sweep-line method

- The basic idea of the **sweep-line method** is to exploit a certain kind of **progress** exhibited by many systems.
- Exploiting **progress** makes it **possible to explore** all the reachable markings of a CPN model, while only storing **small fragments** of the state space in main memory at a time.
- This means that the **peak memory usage** is **significantly reduced**.
- The sweep-line method is aimed at **on-the-fly verification** of **safety properties** (e.g., determining whether a reachable marking exists satisfying a given predicate).

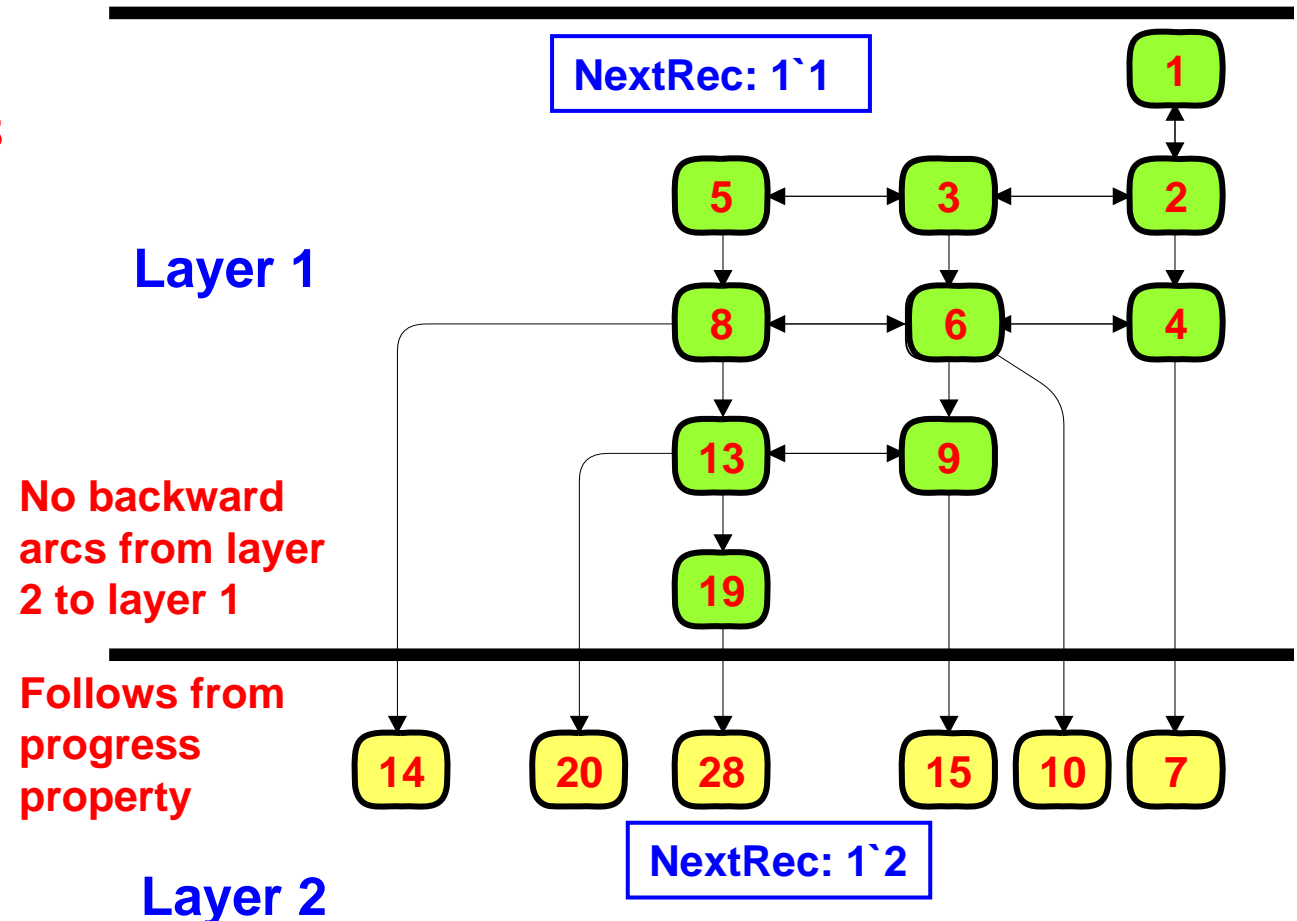


Simple protocol



Initial fragment of state space

- Each marking has **successor markings** either in the **same** layer or in **higher** layers – never in lower layers.



We process markings layer by layer

- We **process the markings** (i.e., calculate successor markings) **one layer at a time**.
- We only **move** from one layer to the next when **all markings** in the first layer have been **processed**.
- We can think of this as a **sweep-line** moving through the state space (layer by layer).
- At any time during state space exploration, the **sweep-line** corresponds to a **single layer**.
 - All **markings** in the layer are “on” the sweep-line.
 - All **new markings** calculated are either **on** the sweep-line or **in front** of the sweep-line (i.e. in a higher layer).



Progress measure

- The **progress** in the protocol system is captured by a **progress measure** which is a function mapping each marking into a **progress value**.

Converts a multi-set 1^x with one element to the colour x

```
fun ProtocolPM n = ms_to_col (Mark.Protocol'NextRec 1 n);
```

- **Monotonic** progress measure:

$$M' \in \mathcal{R}(M) \Rightarrow \text{ProtocolPM } M \leq \text{ProtocolPM } M'$$



Statistics for sweep-line method

Limit	Packets	Nodes	Arcs	Nodes (peak)	Nodes	Time
1	4	33	44	33	1.00	1.00
2	4	293	764	134	2.19	1.00
3	4	1,829	6,860	758	2.41	1.00
4	4	9,025	43,124	4,449	2.03	1.78
5	4	37,477	213,902	20,826	1.80	1.65
6	4	136,107	891,830	82,586	1.65	1.51
4	5	20,016	99,355	8,521	2.35	1.95
4	6	38,885	198,150	14,545	2.67	2.19
4	7	68,720	356,965	22,905	3.00	2.27
4	8	113,121	596,264	33,985	3.33	2.41

Configuration

Standard method

Sweep-line

Gain



UNIVERSITY OF AARHUS

Modelling and Validation of Distributed Systems Group
Department of Computer Science

Kurt Jensen and Lars M. Kristensen
Coloured Petri Nets

14

Summary for sweep-line method

- From the **statistics** on the previous slide, it can be seen that the **sweep-line method** yields a reduction in both **space** and **time**.
- The **space reduction** was expected since markings are deleted during state space exploration.
- The **time reduction** is because the **deletion** of states implies that there are **fewer markings** to compare with when determining whether a marking has been seen before.



Generalised sweep-line method

- Above we have used a **monotonic** progress measure:

$$M' \in \mathcal{R}(M) \Rightarrow \text{ProtocolPM } M \leq \text{ProtocolPM } M'$$

- It is also possible to use a **generalised** sweep-line method where the monotonicity property only is satisfied by most steps.
- The **generalised** sweep-line method performs **multiple sweeps** of the state space, and it makes certain **markings persistent** which means that they cannot be deleted from memory.
- The sweep-line method has also been **generalised** to use **external storage** such that **counter examples** and **diagnostic information** can be obtained.
- This is not possible in the basic method since it **deletes** the markings from memory.



Method 2: Symmetry method

- Many **concurrent systems** possess a certain degree of **symmetry**.
- They may e.g. have **similar components** whose **identities** are **interchangeable** from a verification point of view.
- The basic idea in the symmetry method is to represent **symmetric markings** and **symmetric binding elements** using **equivalence classes**.
 - Each **node** represents a **class of equivalent markings** (instead of a single marking).
 - Each **arc** represents a class of **equivalent binding elements** (instead of a single binding element).

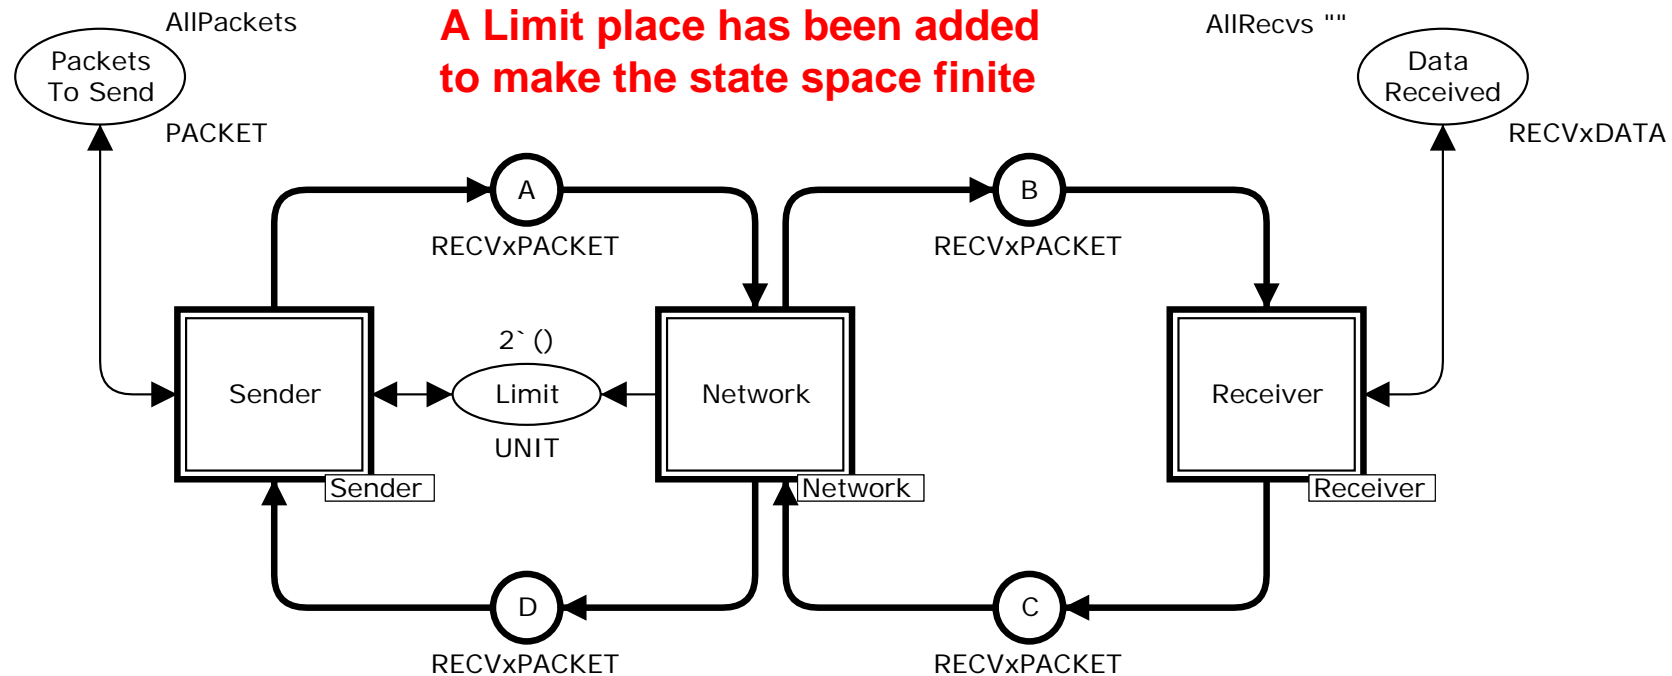


Construction and analysis

- **Symmetry condensed state** spaces are typically **orders of magnitude smaller** than the corresponding full state spaces.
- They can be **constructed** directly **without** first constructing the full state space and then grouping nodes and arcs into equivalence classes.
- Furthermore, **behavioural properties** can be verified directly on the symmetry condensed state space **without unfolding** to the full state space.

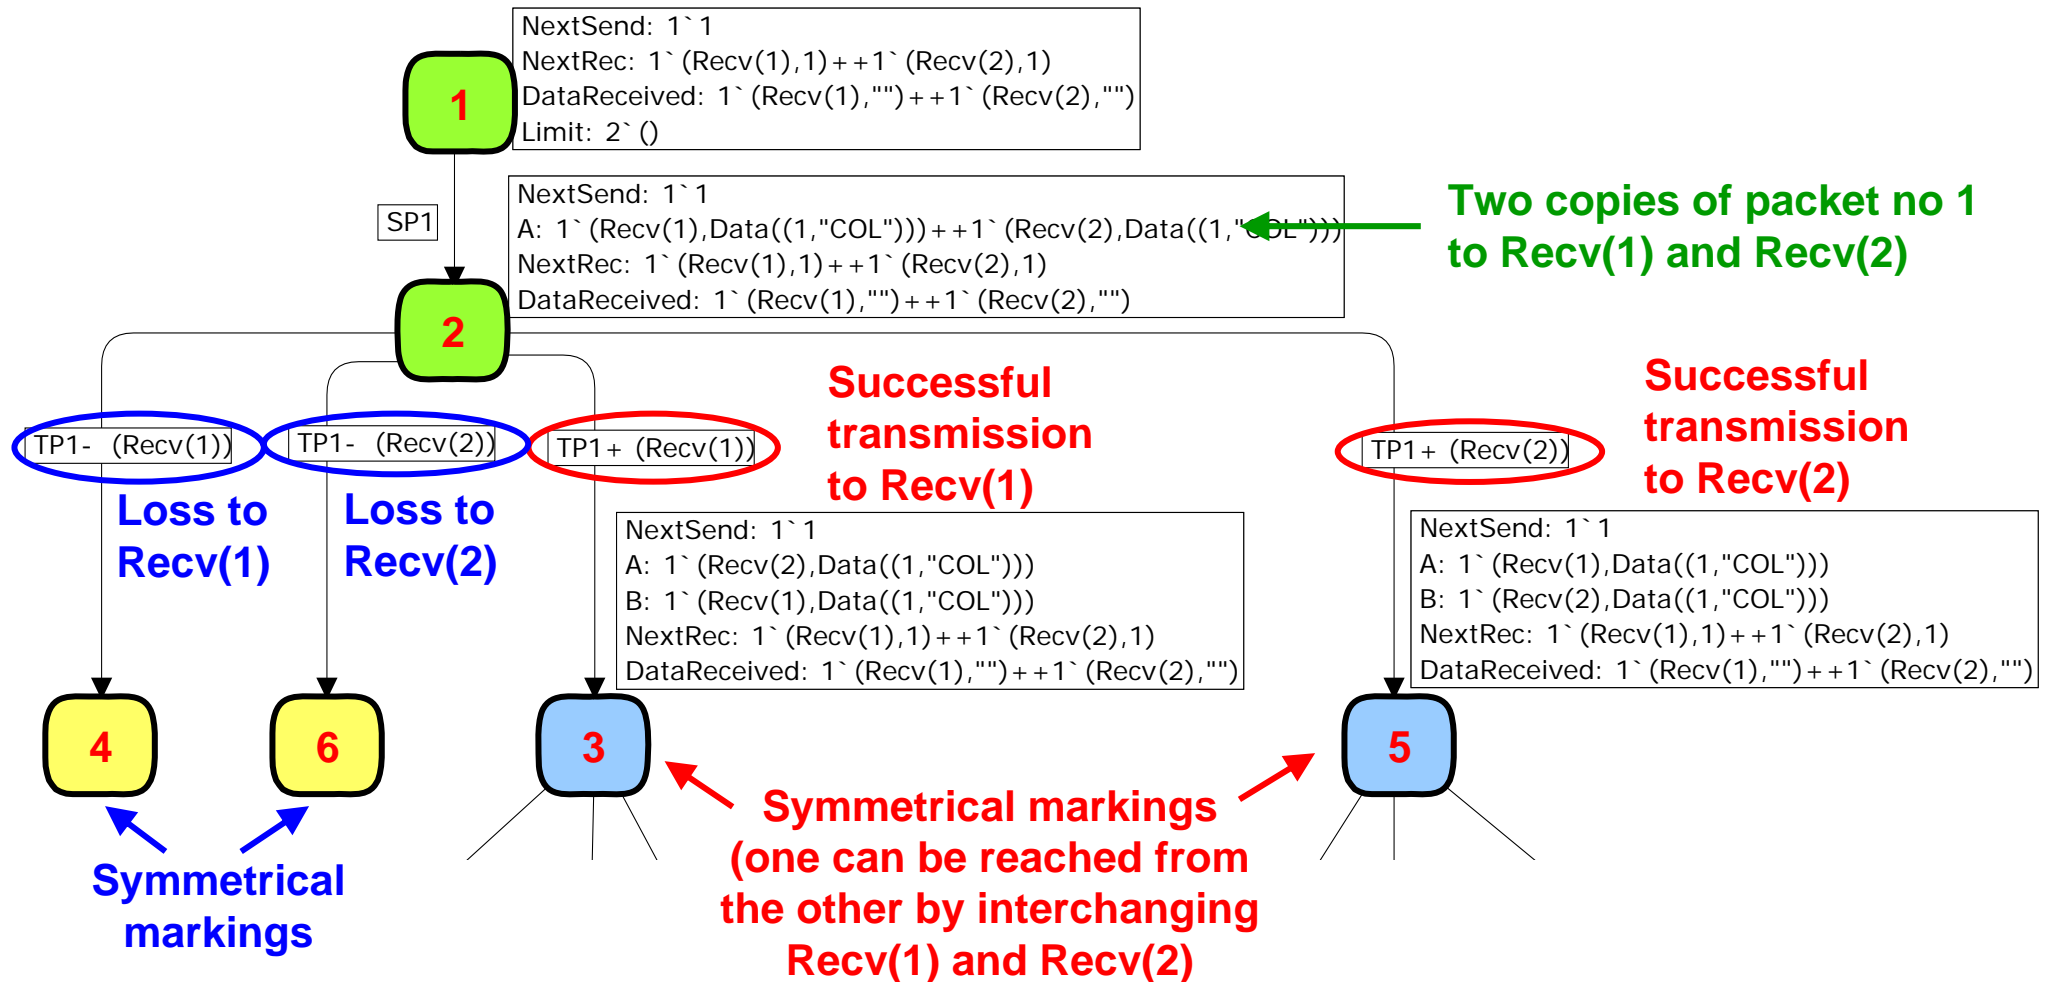


Protocol with multiple receivers

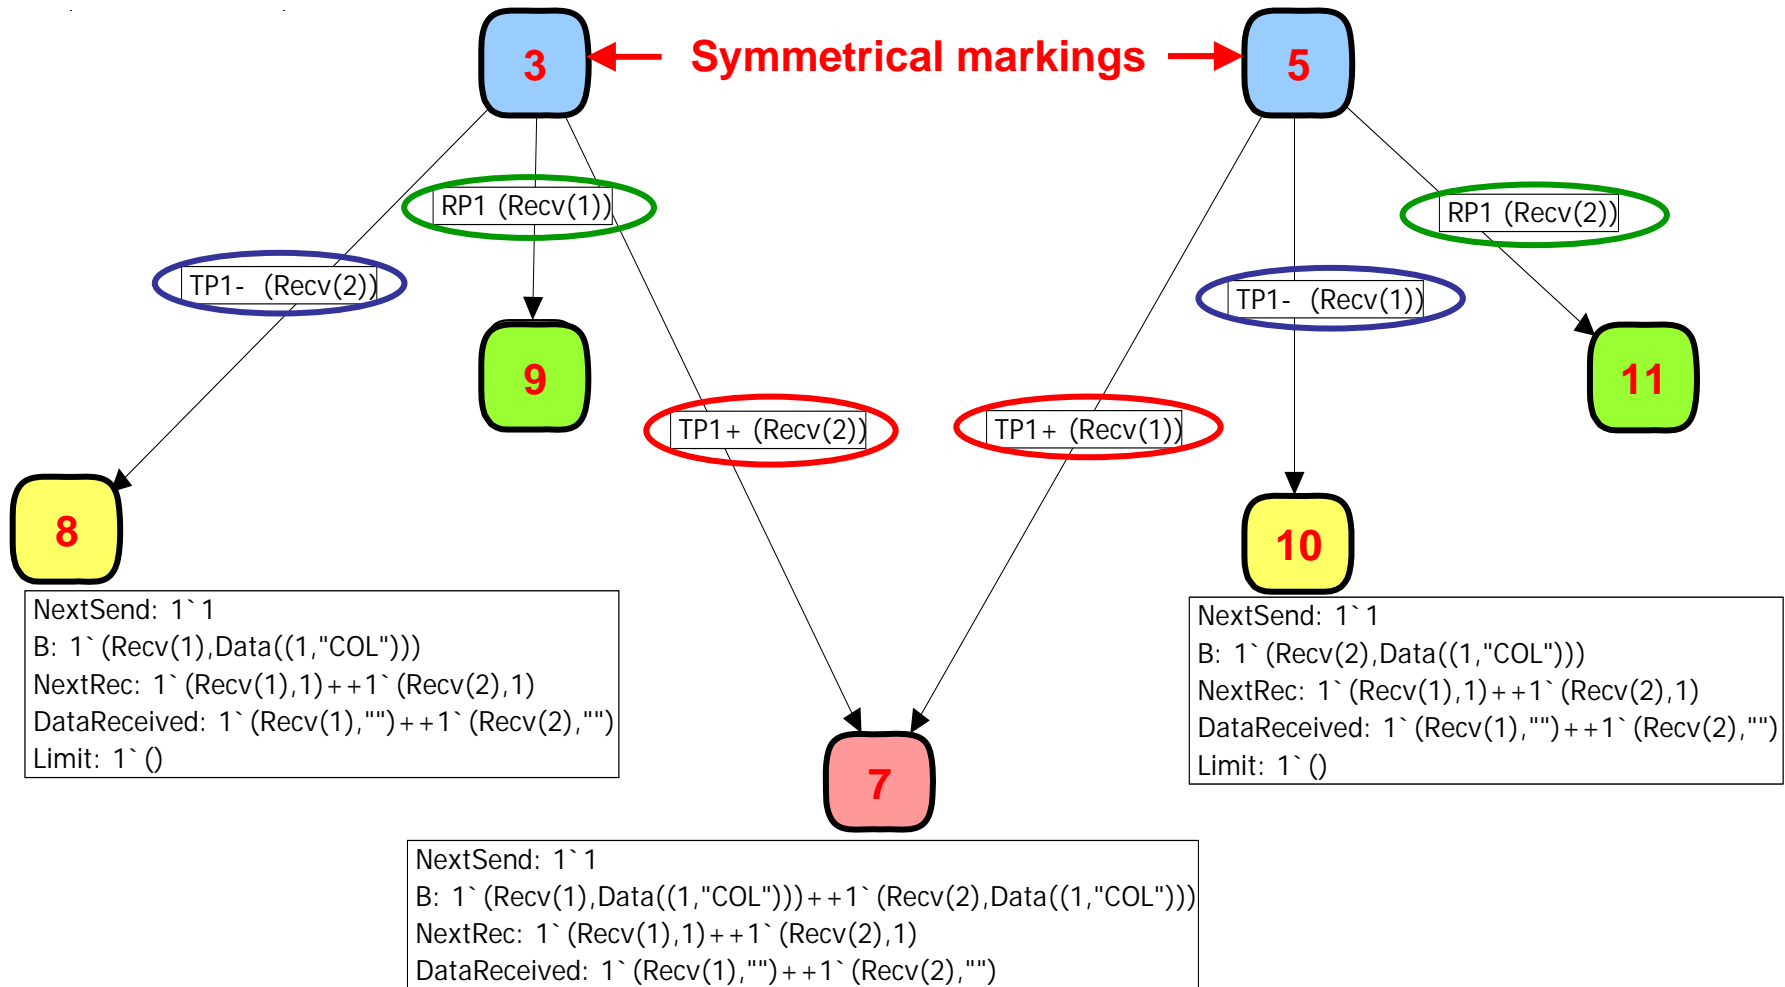


- The **receivers** in the protocol system are **symmetric**, in the sense that they all behave in the same way.
- They are only **distinguishable** by their **identity**.

State space (ordinary)



Symmetrical successors

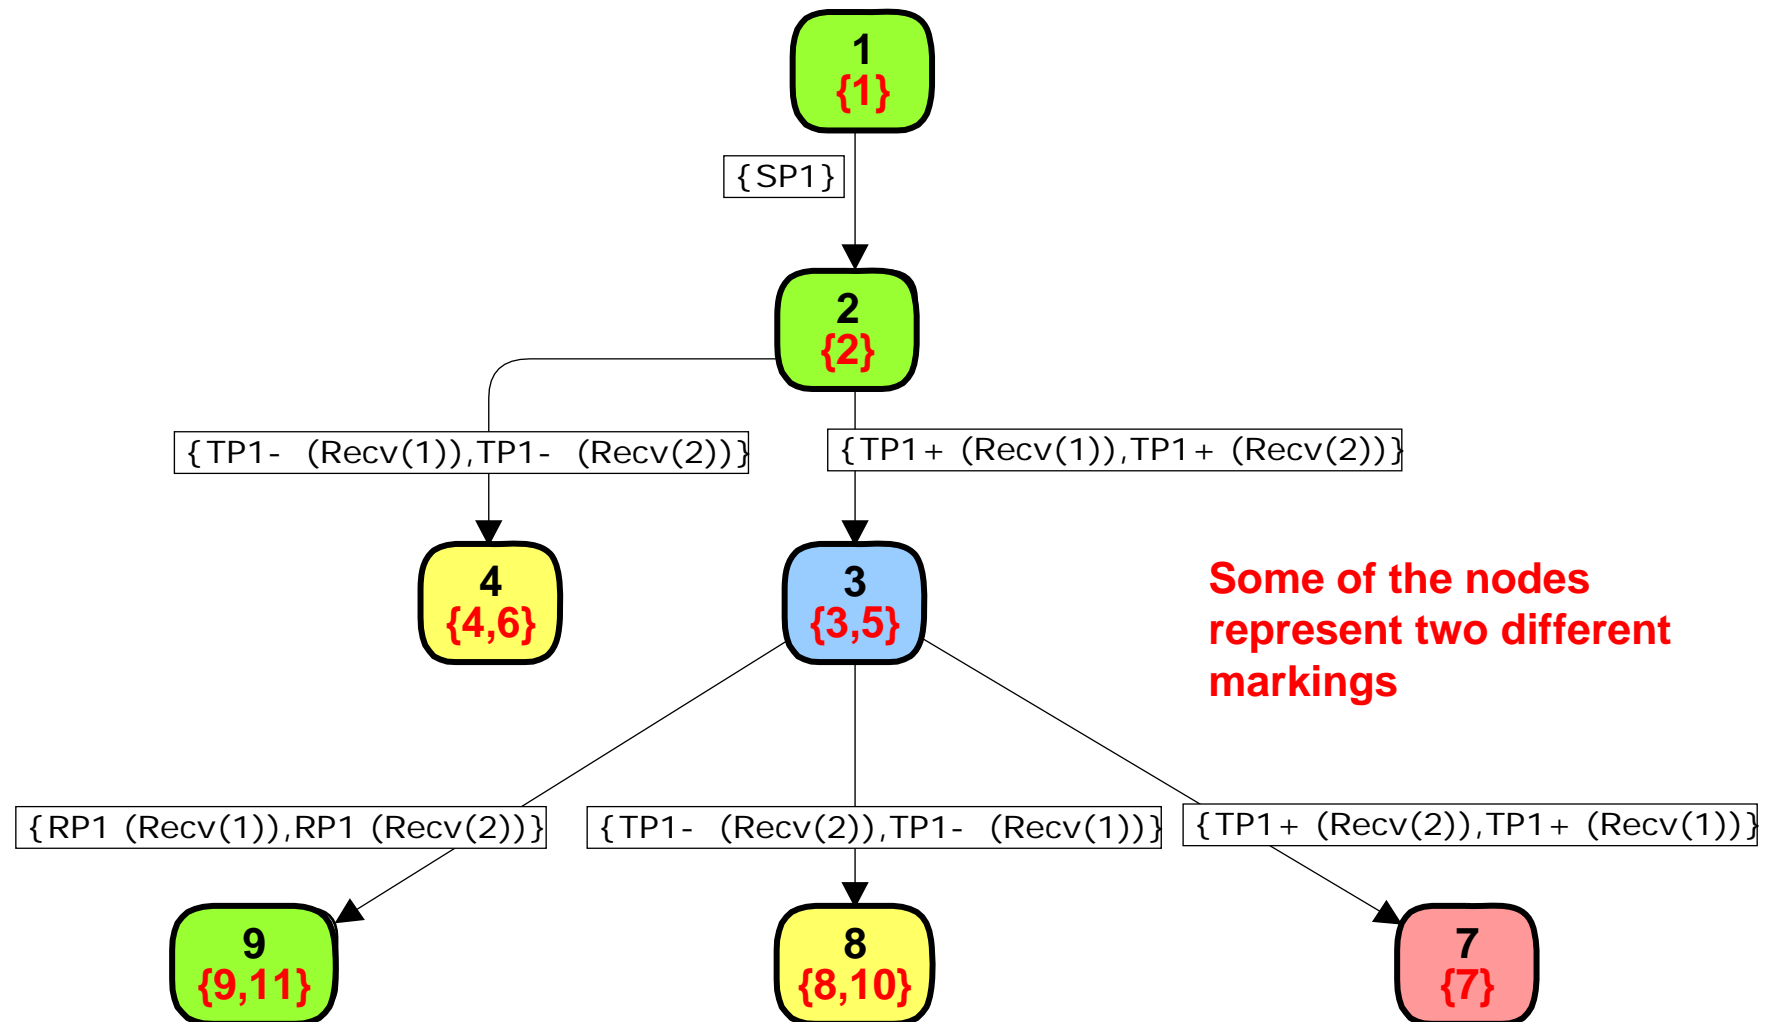


Symmetrical markings

- On the previous slide we saw that the two **symmetric markings** M_3 and M_5 have:
 - **symmetric** sets of enabled **binding elements**,
 - **symmetric** sets of direct **successor markings**.
- By induction this property can be **extended** to finite and infinite **occurrence sequences**:
 - For any occurrence sequence starting in a marking M and all markings M' symmetric with M there exists a **symmetric occurrence sequence** starting in M' .
 - The things which can **happen from M** can also **happen from M'** (up to symmetry).



Symmetry condensed state space



Soundness criteria

- The **symmetries** used to **reduce** the state space are required to be symmetries actually **present** in the CPN model:
- All **initial marking** inscriptions must be symmetric (applying a permutation to the initial marking does not change the initial marking).
- All **guard expressions** must be symmetric (evaluating the guard in a binding must give the same result as first permuting the binding element and then evaluating the guard).
- All **arc expressions** must be symmetric (evaluating the arc expression in a binding and then applying a permutation must give the same result as first permuting the binding element and then evaluating the arc expression).

**Static checks by local
examination of net inscriptions**



Specification of symmetries

- **Colour sets** are divided into:
 - **Atomic** (Int, Bool, String, Unit, enumerations, indexed).
 - **Structured** (products, records, unions, lists, subsets).
- Each **atomic colour set** is associated with an **algebraic group** of **allowed permutations**.
- The **structured colours** sets **inherits** their permutations from the colour sets from which they are constructed.
- **Examples** of permutation groups are:
 - **all permutations** in the colour set,
 - **all rotations** in an ordered colour set,
 - **identity element alone** (no permutation allowed).



Protocol with multiple receivers

- Atomic colour sets:

```
colset NO    = int;
colset DATA = string;
colset RECV  = index Recv with 1..NoRecv;
```

← No permutations (pointing to NO and DATA)
← All permutations (pointing to RECV)

- Structured colour sets:


```
colset NOxDATA = product NO * DATA;
colset PACKET  = union Data : NoxDATA + Ack : NO;
colset RECVxDATA    = product RECV * DATA;
colset RECVxPACKET  = product RECV * PACKET;
colset RECVxNO      = product RECV * NO;
```

← No permutations (pointing to NOxDATA and PACKET)
← All permutations of Recv-component (pointing to RECVxDATA, RECVxPACKET, and RECVxNO)



Statistics for symmetry method

L P R	Nodes	Arcs	Nodes	Arcs	Nodes	Arcs	Time	R!
2 3 2	921	1,832	477	924	1.93	1.98	0.7	2
3 3 3	22,371	64,684	4,195	11,280	5.33	5.73	2.0	6
4 3 4	172,581	671,948	9,888	32,963	17.45	20.38	23.9	24
5 2 5	486,767	2,392,458	8,387	31,110	58.04	76.90	—	120
6 2 6	5,917,145	35,068,448	24,122	101,240	245.30	346.39	—	720



Summary for symmetry method

- **Significant reductions** can be obtained as illustrated on the protocol with multiple receivers.
- The method can be used to check **all behavioural properties** that are invariant under symmetry.
- Computation of the **canonical representations** of markings and binding elements is **computational expensive**.
 - At least as hard as the **graph isomorphism problem** for which no polynomial time algorithm is known.
 - The **present algorithms** exploits a number of advanced algebraic techniques and can efficiently deal with systems where the number of permutation symmetries are **below 10!**
 - This is usually **sufficient** in practice.

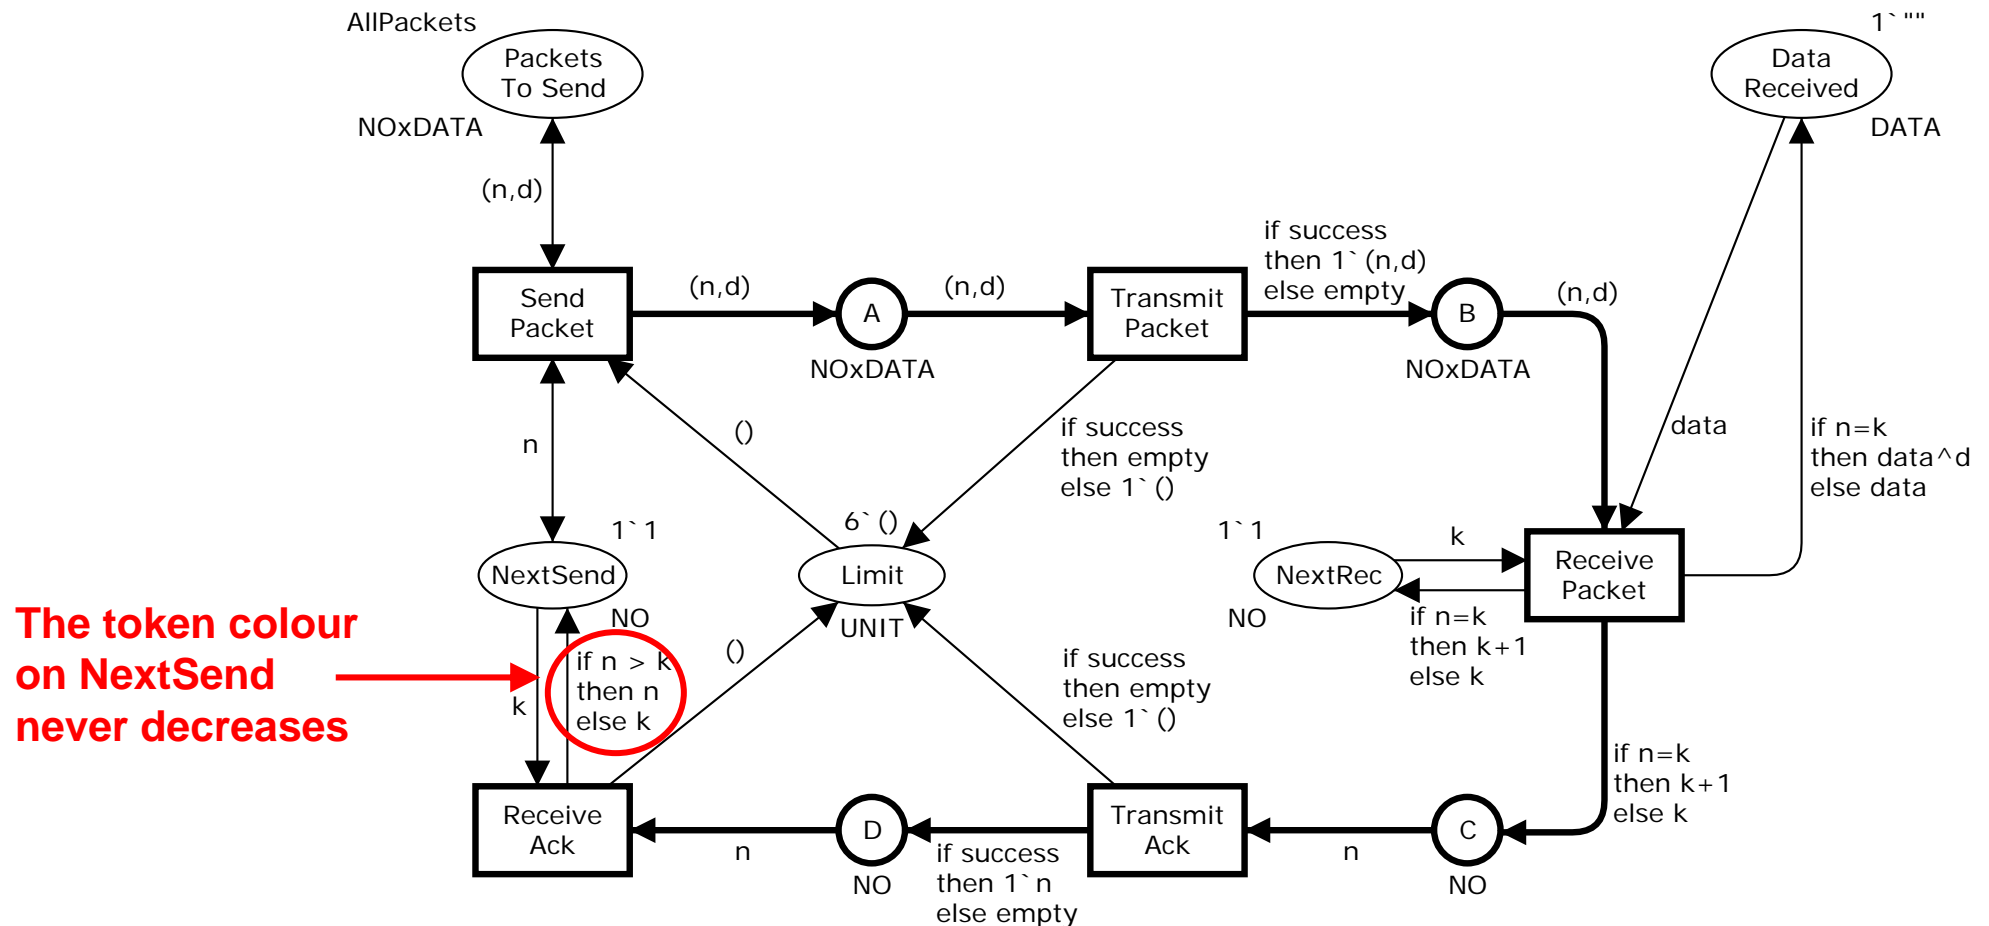


Method 3: Equivalence method

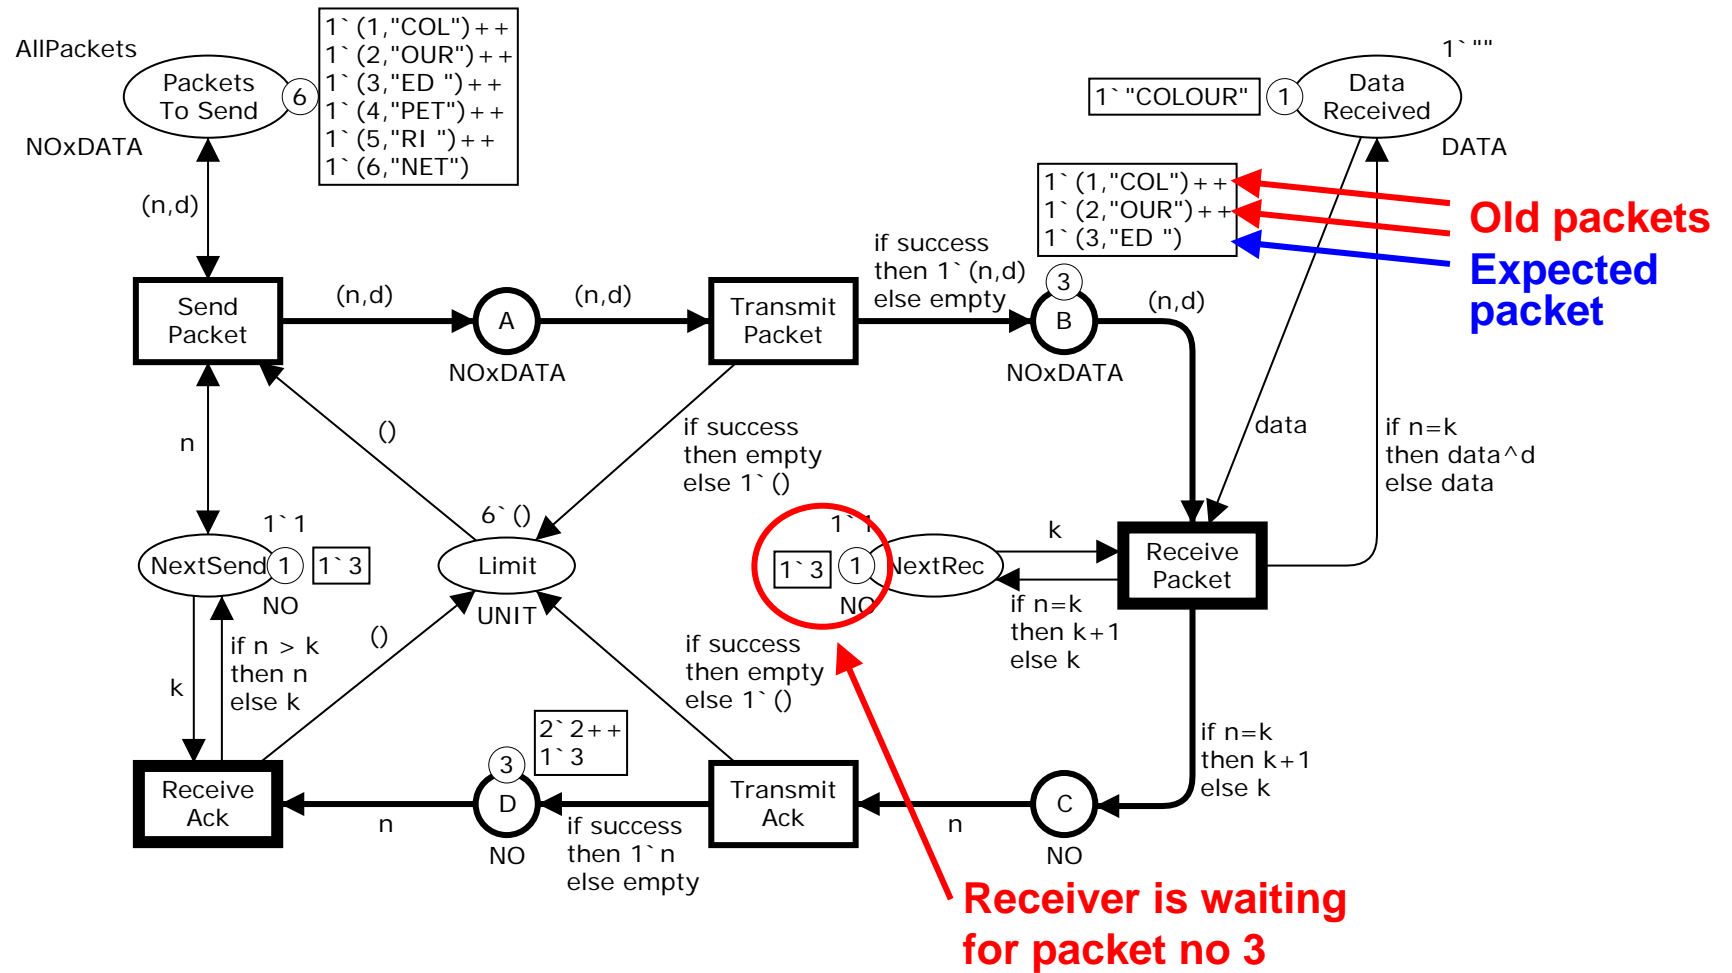
- The **equivalence method** is a **generalisation** of the symmetry method.
- In the **symmetry method** we have **equivalence relations** on the markings and on the binding elements.
- The **equivalence relations** are **induced by the permutation symmetries**.
- In the **equivalence method** the equivalence relations are **specified directly** (without the use of symmetries).
- **Soundness criteria**: Equivalent markings must have equivalent sets of enabled binding elements and equivalent sets of successor markings.



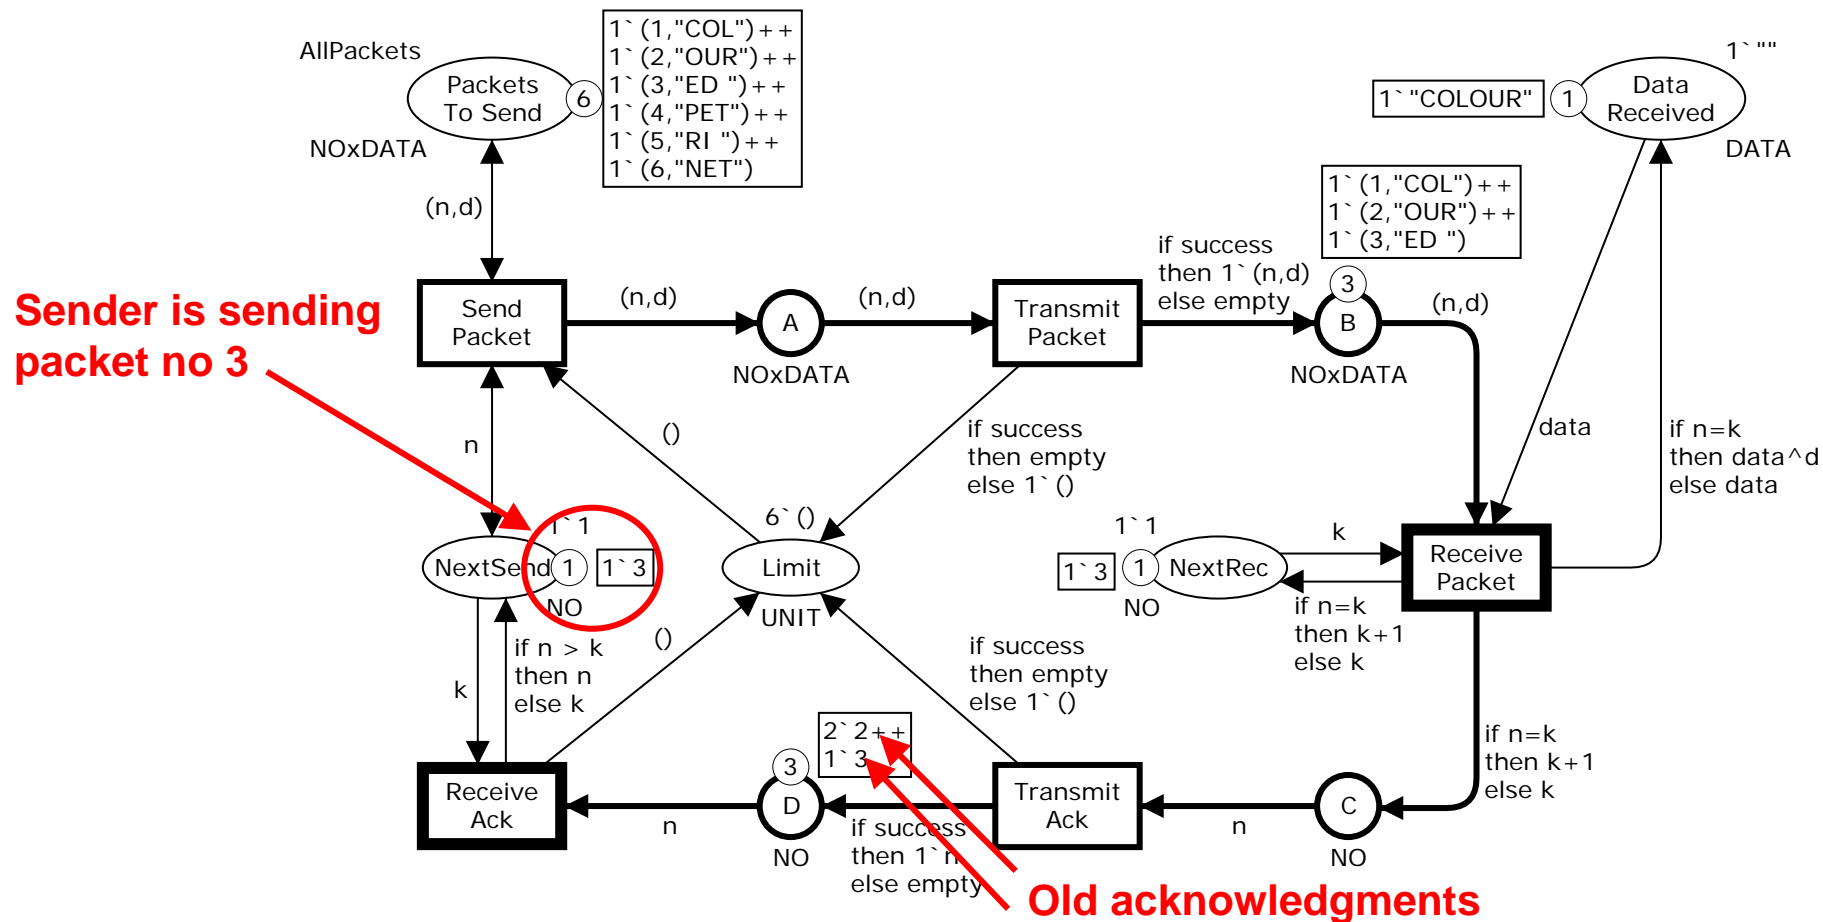
Simple protocol (slightly modified)



Old packets



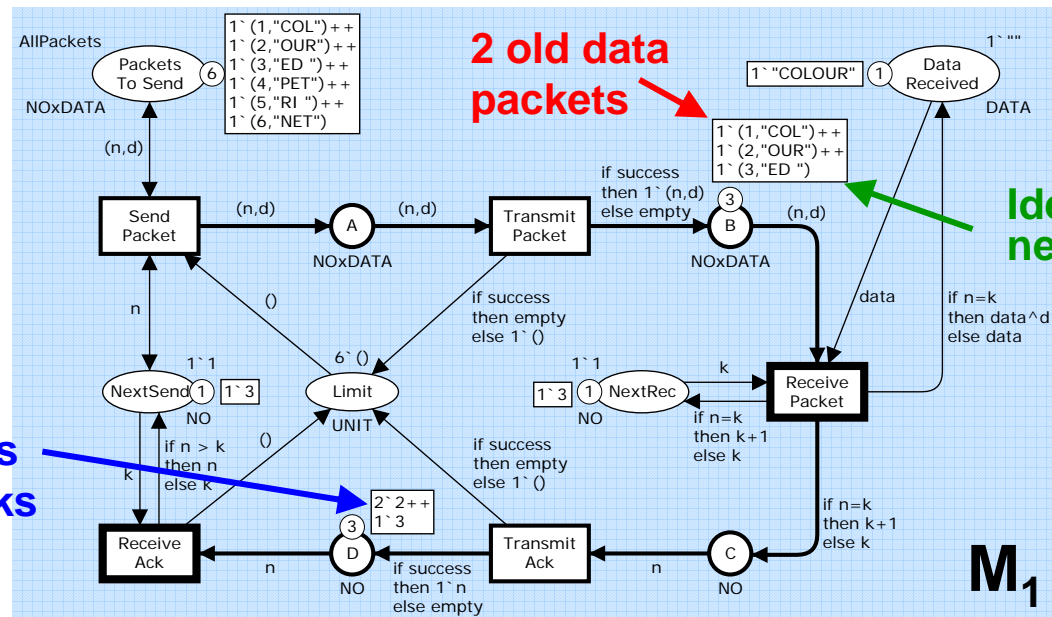
Old acknowledgments



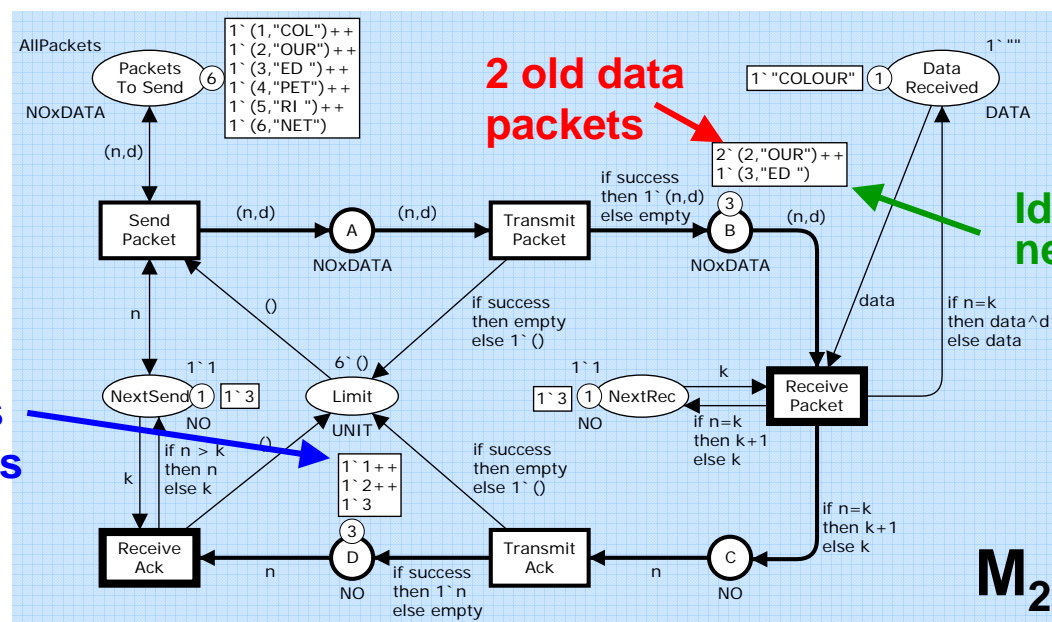
Equivalence relation for markings

- Basic idea:
 - Old data packets can be replaced by other old data packets.
 - Old acknowledgements can be replaced by other old acknowledgements.
- Two markings are equivalent if the following conditions hold:
 - Markings of A, B, C, and D: Identical non-old packets and the same number of old packets.
 - All other places must have identical markings.





The two markings are equivalent to each other



Equivalence relation for binding elements

- Two bindings of the same transition are equivalent to each other if they both involve old data packets or both involve old acknowledgements.
- All other binding elements are non-equivalent.



Statistics for equivalence method

L P	Nodes	Arcs	Nodes	Arcs	Nodes	Arcs	Time
1 4	33	44	33	44	1.00	1.00	1.00
2 4	293	764	155	383	1.89	1.99	1.00
3 4	1,829	6,860	492	1,632	3.72	4.20	0,90
4 4	9,025	43,124	1,260	5,019	7.16	8.59	1.56
5 4	37,477	213,902	2,803	12,685	13.37	18.86	4.09
6 4	136,107	891,830	5,635	28,044	24.15	31.80	13.58

Configuration
Standard method
Equivalence
Gain

L = Limit
 P = Packets



Summary for equivalence method

- The **equivalence method** allows a **more dynamic/general** notion of equivalence than the symmetry method.
- Hence it can be **used** in situations where the symmetry method are of no use.
- The **consistency proof** must be done **manually**.
- The **equivalence relations** must be implemented **manually** (as ML functions).
- Later we shall see that the **equivalence method** can be used to reduce state spaces for **timed CPN models** (without manual consistency proof and with automatic implementation).



Multiple reduction methods

- It is often possible to **simultaneously** use **two or more** state space reduction methods.
- This leads to **more reduction**:
 - in **CPU**, and
 - **memory usage**than each method used in isolation.
- The **sweep-line**, **symmetry**, and **equivalence** methods can be used **simultaneously** with each other.

